Teaching and Learning with Jupyter

Lorena A. Barba, Lecia J. Barker, Douglas S. Blank, Jed Brown, Allen B. Downey, Timothy George, Lindsey J. Heagy, Kyle T. Mandli, Jason K. Moore, David Lippert, Kyle E. Niemeyer, Ryan R. Watkins, Richard H. West, Elizabeth Wickes, Carol Willing, and Michael Zingale

2018-12-02

Contents

Chapter 1

Introduction

This handbook is for any educator teaching a topic that includes data analysis or computation to support the learning—not just courses in engineering or science, but also data journalism, business and quantitative economics, data-based decision sciences and policy, quantitative health sciences, and digital humanities. It aims to give an entry point, and a broad overview of Jupyter in education. Whether you are already using Jupyter to teach, you have found learning materials built on Jupyter that piqued your curiosity, or have never heard of Jupyter, the material in this open book can help you empower your teaching with this new technology.

Project Jupyter is a broad collaboration that develops open-source tools for interactive and exploratory computing. The tools include: IPython, the Jupyter Notebook, JupyterHub, and an ecosystem of extensions contributed by a large community. Jupyter Notebook exploded in popularity since late 2014, fueled by its adoption as the favorite environment for doing data science. It has also grown as a platform to use in the classroom, to develop teaching materials, to share lessons and tutorials, and to create computational stories. Notebooks are documents containing text narratives, combined with executable code (many languages are supported) and the output of that code. This marriage of content and code makes for a powerful new form of data-based communication. Educators everywhere are adopting Jupyter for teaching.

Educators newly adopting Jupyter can be overwhelmed by having to navigate the ecosystem of tools and content. They could study many examples, or consume a myriad of blog posts and talk videos to distill the patterns of good practices and technical solutions to serve their students best. Several early adopters, having much experience to share, decided to begin collecting this know-how, and sharing open documentation about using Jupyter for teaching and learning. The result is this open book: a live document of collective experience with Jupyter in education.

The Jupyter Community Workshop in Washington, DC (November 2018) began that process, with a book sprint aimed at producing the first version of this handbook. The collaboratively written book consolidates explanations and examples covering key topics, including: what is Jupyter, how to try Jupyter, sharing notebooks with students, locally installing Jupyter, cloud offerings, finding example notebooks, writing lessons in Jupyter, making collections for a course, exporting to other formats with nbconvert, writing textbooks with Jupyter, using Binder and JupyterHub, making assignments and auto-grading, making online courses, teaching with Jupyter in the classroom, active learning and flipped learning pedagogies with Jupyter, and guiding learners to create their own content in Jupyter. This open handbook will grow to encompass all you need to know about Jupyter in teaching and learning.

If you find these materials helpful or inspiring, give us a shout-out on Twitter using `#Jupyter4Edu`. We hope you do!

Acknowledgments

GitHub repository for this book: https://github.com/jupyter4edu/jupyter-edu-book

Chapter 2

Why we use Jupyter notebooks

2.1 Why do we use Jupyter?

As teachers we are responsible for many activities, including creating lessons, lectures, courses, assignments, and supportive environments; encouraging engagement and performance in the classroom; helping students learn to think critically so they can become lifelong learners and problem solvers; making material relevant and meaningful to students' diverse interests and backgrounds; assessing student learning (including grading and evaluation); encouraging students to persist with emotional labor (feedback, communication, etc.); and trying out teaching and learning practices that improve our ability to do all of these things.

In short, we design learning environments and experiences.

We use Jupyter notebooks to design learning environments to help support these activities. We believe that incorporating Jupyter notebooks in our teaching has allowed us to improve students' understanding of course content, increase student engagement with material and their participation in class, and to make concepts more meaningful and relevant to students' diverse interests. We represent a variety of disciplines and have many diverse instructional goals, all of which have been supported using Jupyter notebooks. The goal of this handbook is to provide you with ideas to help you address your own instructional and pedagogical goals.

Through a series of anecdotes we will illustrate how you, as an educator, can use Jupyter notebooks to increase your students' 1) engagement, 2) participation, 3) understanding, 4) performance, and 5) preparation for their career. These are starting places and we are confident that you will also take these examples in new and exciting directions.

2.2 But first, what is Jupyter Notebook?

Project Jupyter is three things: a collection of standards, a community, and a set of software tools. Jupyter Notebook, one part of Jupyter, is software that creates a Jupyter notebook. A Jupyter notebook is a document that supports mixing executable code, equations, visualizations, and narrative text. Specifically, Jupyter notebooks allow the user to bring together data, code, and prose, to tell an interactive, computational story. Whether analyzing a corpus of American Literature, creating music and art, or illustrating the engineering concepts behind Digital Signal Processing, the notebooks can combine explanations traditionally found in textbooks with *the interactivity of an application.*

Jupyter is a free, open source platform that is an excellent learning environment for students. For teachers, it increases our efficiency and decreases cognitive load so we can engage students. Notebooks can be useful for achieving your goals as a teacher in numerous environments from STEM labs or humanities narratives, to podium lectures or flipped classrooms. We use Jupyter notebooks in small classes and for classes that

Figure 2.1: A Jupyter notebook, starting with a markdown cell containing a title and an explanation (including an equation rendered with LaTeX). Three code cells produce the final inline plot.

have hundreds of students. Jupyter notebooks can be used for teaching part of one lecture or can be used to teach a whole course. Jupyter notebooks enable us and our students to have a conversation with a problem and link to resources, like audio, video, images, visualizations–and even allow students to mix and remix these. And yet students need to install nothing beyond a modern web browser to use this free software.

Jupyter notebooks can be used to organize classroom materials and objects, store and provide access to reading materials for students, present and share lecture materials, perform live coding, explore and interact with materials, support self-paced learning, grade students' homework, solve homework problems, or make materials reusable to others (see Chapters 3 and 4).

Read on to find out how we have used Jupyter notebooks for teaching and learning to benefit both our students and ourselves. Jupyter notebooks support a wide range of learning goals, including learning to program, learning domain knowledge, and practicing communication skills like storytelling. The authors of this book have used Jupyter notebooks to teach:

- Sciences
  - Physics and astronomy
  - Geoscience
  - Biology
  - Cognitive Science
  - Computer science
  - Data science
  - Statistics
  - Social sciences
- Writing
  - Writing Seminar
  - Writing and technical communication
- Digital Humanities
  - Music
  - Text analysis
  - Metadata processing
- Engineering
  - Chemical engineering (kinetics and reactor design)
  - Mechanical engineering
  - Aerospace engineering
- Introduction to Programming
  - High school
  - College and university-level courses (true introductions through advanced courses)

Our other use of notebooks for education include:

- Building models/simulations (with and without programming)
- Using widgets to demonstrate and interact with simulations
- Visualizations of process and data

2.3   Course benefits & anecdotes

2.3.1   Engagement

As teachers we routinely struggle to engage our students, especially when we are constrained by the format of the course (e.g., online, 50-minute lecture), available technologies, students distractions, and/or other factors. Nevertheless, it is substantially our responsibility to create environments and experiences within these limits that engage students in our courses. This is where notebooks can give you another tool to break out of the mundane, and get students engaged in their learning.

Figure 2.2: From: http://go.gwu.edu/engcomp2lesson4

2.3.1.1 Conversations with data

The creators of Jupyter describe it as a set of open-source tools for interactive and exploratory computing, and a platform for creating computational narratives. Jupyter allows us, as educators, to narrate a "conversation between the student and data". Consider this example, using the data of life expectancy of many countries over the years:

> I use a short bit of code to make a graph showing the time evolution, in what is called a "spaghetti plot" (see figure). Looking at this messy graphic, I point out how most of the lines show growth over time: life expectancy is improving all over the world. But a couple of lines show a marked dip in a given year. I can ask students: which country had that dip? What happened there? Why? With a bit more coding, we identify that Cambodia had a shocking life expectancy of about 30 years in 1977, and Rwanda had even worse life expectancy in 1992. We then have the opportunity to discuss why these countries experienced a mortality crisis. The data brings to life a meaningful discussion, with many possible paths involving history, politics, economics, and health. – Lorena Barba

Jupyter notebooks are essential tools of connection—tools that engage learners in transitions in their thinking. The opportunity of intermingling computation into a narrative, creating a conversation with data is a powerful and effective form of communication. With Jupyter, you now have a new form of content to create and share with learners: *computable content*. In a world where every subject matter can have a data-supported treatment, where computational devices are omnipresent and pervasive, the union of natural language and computation creates compelling communication and learning opportunities.

2.3.2 Participation

Engaging students in your courses requires their participation and interaction with you, their peers, and/or the content (M. G. Moore, 1989). How, when, and why you use student participation in yours will, of course, depend on your goals, the specific objectives for teaching the content within your course, your students, and other factors. Using notebooks, however, encourages participation and gives you more tools for promoting participation. Notebooks can connect students to authentic external audiences as well. Students can, for example, consume notebooks from other classes, and publish notebooks where others can read them.

2.3.2.1 Real world experience – bringing concepts to life

Notebooks are living documents, meaning they can be edited to respond to questions or input from students and used a conversation piece during a lecture or presentation.

> Our group uses Jupyter notebooks as "apps" to demonstrate concepts in geophysics. These notebook-apps connect numerical simulations to widgets and relevant plots. In the classroom, we ask students to help define input parameters based on an application or case study that they are interested in. Prior to displaying the results, we ask students to build a mental image of their expectations. If the resultant image matches their expectations, then we have reinforced a concept, and if not, it is an opportunity to learn. We as instructors can interactively engage with students' questions by updating the inputs to the simulation in order to explore concepts with them. Students have access to the same notebooks through free web-platforms like Binder, so simply by following a link, they can take the steering wheel and engage with the concepts on their own. Notebooks bring the concepts to life and serve as a conversation piece for the interaction between learners and educators. – Lindsey Heagy

2.3.2.2 Real world experience – Ticket to leave

Another example of generating participation in the classroom with Jupyter notebooks is the Activity magic, available as an extension. It creates what has been called a "ticket to leave" (or "exit ticket") via the notebook. The idea of a "ticket to leave" is an excellent way to end a class or lab. Briefly, it is just a survey that you give the students (see figure). Often, these surveys are given via a Personal Response System (also known as "clickers" or PRS) or cell phones. There are a few uses of such surveys:

1. Give the instructor some feedback on the students' understanding, as a whole
2. Provide time and opportunity for students to review and synthesize today's materials
3. Allow the students to apply their recent knowledge to a novel problem
4. An additional instance to learn the materials

These questions do not typically require much time to answer, but are meant to capture the essence of the conversation of the class. After a minute or so to contemplate the question, the students select their answer (by clicking one of the buttons), and instructor shows the gestalt results (see figure).

Good "exit ticket" questions can be domain specific questions, but can also be metacognitive questions (about one's learning style, for example), or high-level organizational questions (e.g., "what was the goal of today's discussion?"). We recommend leaving enough time at the end of class (perhaps 10 minutes) to have a full and complete wrap-up discussion. After the discussion, you may wish to adjust the following class meeting if you feel that not enough students had the insight you were aiming for. For more information on "tickets to leave" see https://www.brown.edu/sheridan/teaching-learning-resources/teaching-resources/course-design/classroom-assessment/entrance-and-exit/sample

Figure 2.3: Dr. Douglas Oldenburg (left) engaging with a student during a short course on geophysical electromagnetics (https://geosci.xyz). Photo credit: Seogi Kang

Figure 2.4: Example of the Activity magic seen from the students view. A question, with multiple choice answers is shown, with buttons for their input.

Figure 2.5: The Activity magic, from the the instructor's perspective. The barchart is shown on the project once all of the students have had a chance to respond.

2.3.3 Increasing understanding

Within any course you will typically try to achieve a diverse set of objectives. Benjamin Bloom (https://en.wikipedia.org/wiki/Bloom%27s_taxonomy) provided a framework for the detailed objectives we want to achieve, ranging from basic knowledge (such as, terminology, specific facts, trends and sequences, classifications and categories, etc.) all the way to ability to evaluate and create (such as, abstract relationships, judgments but based criteria, original works). Achieving the former (i.e., basic knowledge and comprehension) is far easier to achieve that understanding (i.e., evaluation and creation); yet, most often we, as educators, are striving for increasing the complex understanding of our students on the topics we are teaching. The good news is that notebooks offer a valuable tool for teaching toward understanding – moving students, for example, from passively viewing course content to exploring, analyzing, synthesizing, and evaluating the content in active ways.

2.3.3.1 Real world experience – Guiding learners at their own pace

The fundamental theory behind Computational Fluid Dynamics (CFD) used in Aerospace Engineering is based on understanding the Navier-Stokes equations. "CFD Python" is a collection of Jupyter notebooks based on a practical module that Lorena Barba began using in class in her Computational Fluid Dynamics (CFD) course at Boston University in 2009. The 5-week module develops worked examples that build on each other to incrementally guide the learner to create a program to solve the Navier–Stokes equations of fluid dynamics, in 12 steps.

> In 2013, I was invited to teach a 2 day mini-course in the Latin-American School in High-Performance Computing, in Argentina. The Jupyter notebooks platform allowed me to create a guided narrative to support learners with different background experience and knowledge. For that event, we wrote notebooks based on the CFD course module, to use as instructional scaffolding in the minicourse. Twenty students worked through the notebooks as self-paced lessons, while I went from desk to desk asking and answering questions. About four of the students completed all 12 steps in the 2 days, a bulk of them achieved up to about Step 8, and a few of them lagged behind in Steps 4 or 5 by the end of the course. For those who completed the full module, they

had achieved in 2 days what my regular students in the classroom normally took 5 weeks to do. Seeing that was an eye-opening moment: both the power of worked examples in code, and the ability to allow learners to follow their own pace made a remarkable difference in these learners. – Lorena Barba

Based on the experience developing the "CFD Python" learning module (Barba & Forsyth, 2018), this basic design pattern was adopted for creating lessons using computable content:

1. Break it down into small steps
2. Chunk small steps into bigger steps
3. Add narrative and connect
4. Link out to documentation
5. Interleave easy exercises
6. Spice with challenge questions/tasks
7. Publish openly online

This was particularly helpful for student understanding.

2.3.4   Increasing student's performance

The goal of learning is often actualized through the performance of students. This is routinely most visible by what we attempt to assess during and at the end of instruction. Using notebooks we can create a variety of a performance opportunities for students, thereby giving them more opportunities for practice and feedback, as well as more opportunities for us, as instructors, to assessment their ability to perform.

2.3.4.1   Real world experience – The worked-example effect

The worked-example effect is the best known and most widely studied of the cognitive load effects (Sweller, 2006). It refers to providing full guidance on how to solve a problem, resulting in better student performance than problem-solving conditions with no guidance. For complex tasks, inexperienced or beginner learners benefit the most from the worked-examples procedure. One study (Chen, Kalyuga, & Sweller, 2015) concludes that: "worked example effect occurs for complex, high-element interactivity materials that impose a heavy working memory load" and "when dealing with complex material that learners may have difficulty understanding, high levels of guidance are likely to result in enhanced performance over lower levels of guidance." This research-based guidance seems especially relevant for teaching novice programmers to use computation in the context of their subject matter (science, engineering, or other).

2.3.5   Increasing students' preparation for their career

In preparing students to apply what they have learned, striving to align what happens in the course with what they will experience in their career is important. From using parallel software to mirroring workflows, we want our students to experience and be prepared for the workplace. Recognizing, of course, that workplaces are not static and the skills required for a career are always emerging, using notebooks provides a flexible platform to build skills and build portfolios of what students can do.

2.3.5.1   Real world experience – Publishing a data narrative as a demonstration of industry ability

For Data Science careers, a publicly shared narrative about a data analytics project goes a long way at demonstrating the student's potential at an interview. Elizabeth Wicks has her students develop a Jupyter notebook that tells the story of a data munging and analysis project done in the class. The students then

publish this notebook to their Github profile pages. Being that Jupyter is one of the most popular ways in industry to communicate data science results, the students have a very valuable key to a potential career.

TODO: Add quote from Elizabeth

2.4 Student benefits

Creating opportunities for students to develop as learners stretch beyond the boundaries of any specific course where you may use notebooks. By enriching their learning experience in your course, you will help them develop valuable skill-sets and mind-sets that they will take with them into other courses and into their career.

2.4.1 Computational thinking

Jupyter notebooks support a wide range of learning goals. Its interactivity enables building intuitive understanding of domain knowledge, such as the understanding of a mechanical response of a system while varying parameters or understanding how an algorithm behaves. Notebooks can also help teach effective communication skills, combining prose with graphics into a strong narrative. Finally, notebooks can support teaching or strengthening programming skills, by combining code with text descriptions and visualizations. Even if a notebook is designed to be consumed passively, the exposure to code helps show students how to do something—and that they can do it themselves. This also helps demystify coding for students who do not view themselves as traditional "computer science" types.

Using notebooks, you can create rich learning experiences that link together the core foundations of computational thinking:

- *Decomposition*: Breaking down data, processes, or problems into smaller, manageable parts
- *Pattern Recognition*: Observing patterns, trends, and regularities in data
- *Abstraction*: Identifying the general principles that generate these patterns
- *Algorithm Design*: Developing the step by step instructions for solving this and similar problems

2.4.2 Open-source

Integrating notebooks into classes also exposes students to a large and growing ecosystem of open-source tools. This supports their education, but also provides experience in the same environment of tools used in industries in high demand for trained employees, such as data science and machine learning. The open-source nature of these tools also ensures that course content remains accessible and affordable to all students—including those outside the traditional university environment.

Unlike proprietary notebook technologies such as Mathematica, or specific programming languages/environments such as Matlab or C++, the barriers to entry for students learning with Jupyter notebooks can be extremely low. At a minimum, during a lecture, students can simply watch/read an interactive demo using a notebook, to replace slides or lecture notes. On their own, using a cloud service such as Binder or JupyterHub, students can open any modern web browser to some address and interact with a notebook (an example of this technology can be found at https://jupyter.org/try), without needing any installation or configuration. In the most complicated case, students can install Anaconda and follow simple instructions to install the Jupyter Notebook, which works and looks the same on all platforms—and is free and open source.

2.4.3 Active learning

Thanks to their interactivity, notebooks enable a spectrum of active learning methods, which have been shown to increase performance in science, engineering, and mathematics (Freeman et al., 2014). To start, students

can consume notebook content by reading and running notebooks, then move to editing or completing notebooks as assignments. This allows students to focus on the content and concepts, rather than just note-taking.

At the top of Bloom's Taxonomy is pure creation, where students can, for example, author complete computational essays. In both cases, notebooks support courses where students have a wide range of experience and ability: students who need help can rely on the scaffolding of prose explanations and existing code, while also providing room to stretch and explore for more-experienced students. The additional annotation and prose that accompanies code also helps support non-traditional learners and students from underrepresented groups who may have less initial experience/comfort with programming.

Instilling the habits of active learning, through the use of notebooks, will also provide benefits beyond the boundaries of your course. Interactivity drives engagement, interest, and exploration of concepts. Engaged students in your course, are more likely to be engaged learner in other courses and beyond.

2.5   Instructor benefits

Notebooks can be adopted at a variety of levels and formats, offering flexibility based on the needs of a course and comfort/interest level of the instructor: in-class demos, interactive labs, auxiliary material (e.g., book replacements, lecture note supplements), assignments, or full course content in a flipped learning environment. Notebooks offer a route to active learning methods for instructors without experience of them, but do not force a particular teaching style.

At a minimum, notebooks can be used to make publishable and interactive lecture notes that blend narrative text, images, videos with image and results to present the concepts. Furthermore, these course materials can be developed gradually, starting with a low-effort draft to a more-polished, publishable document that can be easily extended over time—and adopted by others. The growth of open-source communities around software tools and educational resources creates more opportunities for the re-use and adaptation of existing resources.

While many notebook authors do use Python, the Jupyter Notebook supports many languages, so students (and instructors) are not tied to one specific language. Indeed, the name Jupyter comes from three languages: Julia, Python, and R. Furthermore, these (free) tools have minimal barriers to entry—using a cloud infrastructure means students and instructors do not have to install anything, while in the "worst" case installations require a few command-line excursions, but these are free, openly available, and cross-platform.

2.6   Conclusions

We hope that this chapter has illustrated that teaching with Jupyter notebooks can be valuable for you and your student. We have notebooks to be a tool that can increase student engagement, participation, understanding, performance, and preparation for their careers. These are substantial accomplished that can be achieved in a variety of disciplines and content areas. Using several real world examples, we attempted to illustrate the numerous ways teachers are using notebooks. Hopefully these, when combined with the chapters that follow, will guide you in 1) supporting your students' learning, 2) giving you confidence that you can use notebooks, 3) help you understand the necessary logistics, and 4) help give you clear expectations of what can be accomplished with Jupyter notebooks.

Chapter 3

Notebooks in teaching and learning

Jupyter notebooks are a valuable tool for teachers, but their value can only be leveraged if you apply them correctly within the context of your course. In this chapter, you will learn how teachers can initially structure the design of their course and then determine when and how notebooks can be used to achieve their goals.

3.1 Oh the places your notebooks will go!

3.1.1 Introduction

In Chapter 4 you will learn about the many creative ways that notebooks can be used within the design of your course. Many times notebooks can be adapted into course activities that you are already doing, and other times notebooks will give you opportunities to extend what you have done in past to increase the engagement, participation, understanding, and performance of your students. Jupyter notebooks can be a valuable member of your existing instruction toolkit, useful at every point in the learning environment. New adopters of Jupyter can start small, incorporating notebooks into single modules, assignments, or classroom activities. This is an excellent approach to see how your learning audience interacts with the notebook environment and explore notebook hosting systems in a low cost/risk way.

Instructors adopt Jupyter within their classrooms at a variety of levels, each making use of the strong features of the environment. Transitioning an existing course to any new platform seems daunting, but Jupyter notebooks are modular and ideal for an iterative development approach of adoption. Some may find themselves inheriting a course already built in Jupyter, while others will choose to build new courses entirely within Jupyter.

This section seeks to inspire you about the many uses of Jupyter within classroom content and presentation design, and preview how other elements of the Jupyter ecosystem can be integrated into that use. Some of these uses are quick to spin up as experiments to test the waters.

3.1.2 Jupyter notebooks as textbooks

Instructors often write Jupyter notebooks as linear narrative documents. These notebooks are to be read by students and learners, perhaps worked through, marked up and are a relatively one sided information consumption experience.

Most often these notebooks exist as the readings a learner is required to do before class, as reference material (e.g., to review for future assessments), or something that a learner works through on their own as a part

of self study. Jupyter notebooks can be constructed of completely static text, which can serve as a starting point for transitioning existing material into the notebook environment.

This traditional static textbook chapter or section can be extended into an active space by changing inline code text to executable code cells that support modification and experimentation. Adding prompts with suggestions for interrogating the code and examples further extend active learning opportunities without rewriting the original content. Interactive sliders, user input sources, and manipulable visualizations are examples of how other widgets and plugins can open up more possibilities.

As you'll see in later chapters, many authors are using Jupyter notebooks as their primary authoring and publishing platforms. These materials are published on paper and online, meaning that the interactive portions of the learning experience are first class elements within the resource.

3.1.3   Notebooks as workbooks/primer

Workbooks engage students in the notebook environment by including active elements where they are asked to manipulate or create new content. This moves students from a passive or static learning environment like a book into an active learning experience where they can engage with and critically think about the content.

You can include many pedagogical patterns (discussed in the next chapter) within a workbook, crafting a completely custom learning experience. These workbooks can be assigned as independent student learning (for example, pre-work for flipped classroom), or as part of an in class activity for individuals or small groups.

When teaching in the technical space, the exploration of concepts in their context is an ideal environment for showcasing an authentic experience to students. Studying a concept in isolation may reduce the complexity of the problem to be more accessible, but this removes it from the context for why it exists and may make it harder for students to put all the pieces together and synthesize it into their larger problem solving repertoire.

Retention of this context adds the complexity of expository text, technical reference in disconnected documents, or some boilerplate code. Providing this in a large script or lengthy lists of direction can be overwhelming for students to work with. Jupyter notebooks are excellent tools for teaching these complex workflows, because instead of a script with lengthy code comments or having disconnected documentation, this guiding text can be embedded exactly where it is needed where the coding is happening. This guidance material can also be formatted as markdown cells, which are visually distinct from code, making it easier to visually scan and opens the possibility to add helpful formatting.

A variety of activities can be supported with the executable code cells so learners can explore the space in an interactive and iterative environment. They can see and inspect portions of the surrounding code, but aren't required to touch it, maintaining appropriate granularity for assignments and challenges. Using markdown formatting separating sections helps scaffold larger problems and help them really experience a real world workflow rather than statically read about it.

3.1.4   Notebooks as worksheets/drill sets

Many programming and domain tasks have specific syntaxes to be learned and effectively memorized before they can be internalized. Akin to math homework sheets, short worksheets where learners focus on highly granular or decontextualized problem sets lets them practice the complex syntax and procedures in an isolated and highly focused way. A set of small targeted tasks where these complexities can be practiced without the worry of additional syntax errors or other problem solving requirements can reduce the cognitive overhead a student might be overwhelmed by.

The cell based nature of the Jupyter notebook makes interactive code-based worksheets a clean experience for students to run. Each problem prompt can be written in a markdown cell, perhaps referencing an object or data file established in an initial cell. For example, a list or other data structure would be defined at the top and the exercises below are focused on the relevant methods and usage syntax. Each answer would then

be completed by the student in a single code cell just below that markdown cell. This means that outputs and errors stay with the code producing them, so successes or bugs are easily traceable to the source.

Usages of autograding tools or unit tests, as discussed in later chapters, can be added to give students instant feedback about their work. Example or desired outputs could also be reproduced in markdown cells with the question for further guidance.

3.1.5 Notebooks as notepaper or course packets

The ability for a notebook to represent a linear experience with human prose and working code means that these can be used a student's notepaper in class. They can capture the linear narrative structure of a lesson or lecture, and actually run the code they are taking note of. This ensures that what they have written down actually works, and makes for a strongly reusable document for them moving forward with homework.

Encouraging students to use Jupyter notebooks for taking class notes opens up further opportunities to provide scaffolding and support within the classroom. Providing notebooks with outlines of lecture content or other materials covered in class can become an useful active learning and engagement strategy. You might provide a mostly blank notebook with just topic headings, and ask them to take notes in their own style within those spaces. Or you could choose to embed reference notes, examples, and even small activities within the notebook, asking them to both take notes and work through examples with you.

Caution should be taken to encourage students to carefully maintain a linear order of their code in the notebooks. A later chapter has more information on Jupyter specific caveats that students should remain aware of.

3.1.6 Notebooks as an app

Notebooks even have a place in non-coding classroom content or activity. Interactive user inputs like mouse or touchscreen controlled sliders, buttons, highlighting, etc. allow a notebook user to manipulate input parameters for a visualization, tool, or model without directly editing any of the code within the notebook. These strategies support interactive computational exploration, or transform the notebook into an advanced calculator tool for students to use within their homework. These notebooks are then treated as applications that are distributed or made available for students to use during class or explore on their own.

This allows you to make an existing research workflow accessible to novice students as part of a computational module of a foundational class, or mock up content from a static textbook or reading into something for an active learning activity. The adaptability and reuse aspects of Jupyter notebooks also create opportunities for students to take this code further and adapt it for assignments or other research.

3.1.7 Notebooks as lab reports or assignments

There are a variety of assignment deliverables that programming and technical courses may require. Students may be asked to produce essays, presentations, working code, analytics, and even art or music. Many of these deliverables are directly supported within the notebook environment. Any written work could be completed within the notebook environment with markdown, which is ideal for communication content that is driven by data or incorporating code content. For example, a student could write a computational essay within a notebook, and use one of the presentation tools to present a report out in class, all using the same notebook.

Coding assignments can be submitted to a Jupyter supported learning management system and be auto-graded (discussed in later chapters), providing instant feedback and automating the grading process. This opens up self paced and highly scaled options for many courses, particularly open access MOOCs or large sections. Meanwhile, the inline visualization options mean that as assignment with graphical output can be self contained without trying to embed images within a word processing document or attaching a collection of images along with a script.

The multiple conversion and hosting options available to notebooks means that they can be shared or submitted across many formats. For example, conversion to HTML means that there is zero overhead for viewing the content, and support for markdown and PDF opens up accessibility and other publishing platforms.

3.1.8 Notebooks as interactive multimedia platforms

A variety of media formats can be embedded within a notebook, and other tools more offer platforms to more directly connect notebooks with multimedia content. Instruction content might be split up between short videos (often for flipped classrooms) or a variety of static images might be important for an assignment. The markdown cells within the Jupyter notebook provide several ways to place hyperlinks and embed a variety of media.

Several widgets are also available for embedding playable audio and video content (including from streaming video services) directly within the notebook. This creates a cohesive platform experience for the student, so they don't have to exit out or change screens to work on their assignment and reference that content.

Other tools are available for more directly connecting the notebook content to a video guide. Video lectures or content in your courses can get lengthy in running time and associated notebooks are often just as long. Tools like Oriole provide a platform where you can integrate video timestamps into notebooks to create interactive video experiences. You can include, for example, Youtube videos within notebooks, with text and/or coding opportunities before/after the video. Using timecodes, you can also guide students through the videos in tandem with the notebook text. This further extends the ability to create a single cohesive interactive experience without the need for students to go back and forth between materials.

3.1.9 Notebooks as a demonstration platform

Eventually you will have to display a notebook in class. This may be as a demonstration of how to use a notebook, presenting more a traditional style lecture, creating and editing code, or using an interactive feature to explore an experiment. Normal standards for font sizes, organization, and accessibility stand for these cases.

Displaying actual notebook within your presentation is a natural starting point. This content may include text from markdown and LaTeX, code, and independent figures and sketches. For instance the lecturer could display the notebook, slowly scrolling through the material and interacting with cells containing code while also making use of either a digital or analog (e.g. chalkboard) sketching device. Custom styling plugins are available to change the background color, font, and other viewing aspects of the notebook for better presentation quality and accessibility.

Several slide show tools are available, which allow you to markup notebook cell content for a more traditional slideshow presentation mode without having to exit from your standard notebook. These slides can be scattered across notebook content into specific cells that will correspond to individual slides, with the other content ignored from the presentation. This then can be shown alongside the usual notebook interface and can be flipped between the other forms of content.

It is natural for students to read and interact with notebooks in their standard form, using either Jupyter Notebook or JupyterLab, and this can also be used for presentation. Jupyter Lab offers the convenience and uniformity of being able to open and edit source files (`.py`, etc.) within the same environment and without OS- or browser-specific clutter. Full screening the browser is recommended if presenting using this mode. Presentation styles vary widely and span a spectrum from full-prepared notebooks to blank-slate live coding.

Many find it useful to incorporate an alternate modality such as a physical or digital "board" for free-form diagramming, working through a mathematical derivation, or other written procedural task. Notebooks can be a portion of these presentations or the complete environment, depending on your personal instruction style and content needs.

Care should be taken about how the notebook is presented and demonstrated. Doing a live demo gives you a few options. You might choose to:

- Scroll through a notebook
- Step through a notebook by executing the cells in order
- Fill out details or values into a mostly complete notebook
- Tweak or flesh out a notebook with some content
- Add content to a completely blank notebook

Each of these strategies have a place within a classroom, and their use should be informed by audience needs and learning goals. For example, having a notebook with prepared challenges for the end of each module or section, but with blank cells for the content gives you the opportunity to develop code live within class, but within a structure that keeps your workflow organized, and your formative assessments coded directly into your presentation. Incorporating reference information can make these documents more complete for students and answer common questions. There are many possibilities about what you can put in. Keep in mind that students will often ask for you to share copies of these notebooks after the class session is completed.

How much you focus on live coding will likely be determined by the domain content of the class. Programming courses would clearly have a priority for students to have more thorough practice with writing and typing in the code. However, a conceptual class looking at computational models may interactively tweak parameters for a model and discuss only what is happening within that mathematical model. No code is written or directly changed in the process.

3.1.10 Notebooks as a live coding environment

Live coding, as the name implies, involves the active writing of code within the instruction process. This might be part of a recorded screencast or an in person classroom. The process of live coding has several benefits for the student and instructor.

Showing the process of building up code examples showcases the natural non-linear process of how code is crafted, but walking through the logic of a code example slows instruction down and highlights the reason for inclusion of each element within the code.

Introduction of bugs (either purposeful or accidental) to the code has the added benefit of giving the presenter an opportunity to work through the debugging process and demonstrating that perfect code is never created on the first go.

Live-coding can also be an opportunity to provide an active learning experience by providing notebooks with code that has not been completed before the lecture and having students attempt to fill in the missing lines before doing the live-coding demonstration. Feedback on where students are in this process can be a useful way to also judge what students are retaining and are struggling with leading to just-in-time teaching opportunities.

Formative assessment and prediction prompts can also be incorporated either directly into the notebook or as part of the narration of the lecture. Creating live-coding opportunities can be done anytime where a block of code exists but picking out particularly illustrative examples or key points and appropriately scaffolding the example can be critical. For example, if the critical concept is inside of a for loop then only coding the inner part of the for loop can be helpful and not overwhelm the presentation with scaffolding such as the setup. The reverse can also be true however. If the complexity or amount of scaffolding is displayed learners may struggle to understand the scaffolding rather than concentrating on the key concept.

Many instructors utilizing live coding will choose to have students code along with them. This allows them to practice what they are learning, see it in the natural context of their environment, make all the normal mistakes and typos, and all within an environment where they can ask questions (or pause a video). Actively coding along with the instructor also includes a requirement that they are actively listening to the instructor and engaged with the content.

Presentation styles of scrolling or shift + enter can are not live coding, but live demonstrations. While these limit or negate the benefits of the live coding environment, the benefits of speeding up the presentation or running through code that's irrelevant to the learning goals may be more important. Removing the student's active engagement with the content may eventually lead to their disengagement with the lesson, or missing large chunks of information. Instructors should balance their inclusion of live coding and live demonstration to ensure that students are active and engaged with the most important aspects of the lesson.

Information bandwidth in the classroom during a live coding session needs to be carefully managed, particularly when students are trying follow along. The rhythm of live coding roughly has three stages: preparation, typing, and explanation. These three follow quickly in succession but are independent phases. Preparation is the first, where you stop and explain what you are about to do. Typing is the next phase where you should speak as you type but only say what you are typing. This ensures that what the learners are seeing on the screen and hearing from you match. They will likely be looking back and forth between their screen and yours that they often won't be able to stop and follow what you are saying while they are typing. The final stage is to stop and explain what you have typed and what has or will happen when you run the code. You may choose to execute the code and explain the results or include a formative assessment or prediction question before running the code. Pausing to explain the code you just wrote and walk through the results gives students time to catch up to your typing, time to consider what has happened, and a natural place to ask (and for you to as for) questions about what has happened.

Live coding does take practice to get used to, but can be extremely powerful for you to restrain your pace to your learners and to retain engagement with your students.

3.1.11   Conclusion about places

As you have just seen, notebooks provide a flexible tool that can be used in numerous ways to achieve your course goals. Notebooks are flexible enough that you can use them from relatively passive to very active student learning, you can use them in your lecture or in a flipped-classroom environment. There is no single best way to use notebooks in your courses, and you explore the various options you will want to start filling your use of notebooks with a variety of the pedagogical patterns described in the next chapter.

3.2   Before You Begin...

This chapter focuses on course considerations when incorporating notebooks into a class. Experienced instructors may choose to skim parts of this chapter and focus specifically on how Jupyter notebooks would change their current teaching style.

Before you begin adding Jupyter notebooks to your course, take some time to:

- Identify your teaching goals
- Understand your students
- Develop your content strategy
- Consider the context of the learning environment

The Jupyter notebook is a tool; its use in this context is subject to the expectations of the instruction. Setting expectations for learning depends on your goals, your students, the content, the learning environment context, and you.

3.2.1   Identify Your Teaching Goals

As with the creation of a building or robot or book, it is important to begin the process with a clear goal of what you are trying to create and why. The "and why" could be the most important decision you make in

the whole process and it will (or at least it should) guide all the decisions that follow. The why here is not about why to use Jupyter notebooks, but the why is the goal you have for your students.

Is the goal to teach them critical thinking skills, or how to execute a specific set of functions to solve a problem? Is the goal that they will be able to translate mathematical concepts into real world application, or is the goal to teach them how to code? Be specific and clear, and then let the answers guide your decisions.

3.2.2 Understand Your Students

Your students are central to decisions about when and how to use notebooks alongside other tools in your instruction. An obvious illustration of these decisions would be selections made if you compare teaching 5th graders in relation to teaching graduate students. Yet, there are more subtle differences that you will also want to be aware of and monitor. For example, within a classroom, you will have variation among the backgrounds and skill levels of the students. Depending on the domain, some students may have extensive experience with coding in multiple languages while others may be doing their first computational explorations. Keeping students of all different experience levels engaged and excited is challenging.

3.2.2.1 Learn About Your Students

In many circumstances you will have little background information on your students prior to the first day of a class or workshop. You will have to learn quickly and be prepared to adjust the instruction to fit the students in attendance—not the students you wish you had or anticipated having. Some key considerations to learn about your students:

- **Motivation:** Why are your students participating? What are their goals beyond the workshop or class? Will they be applying what they learn soon, or not for months or potentially years?
- **Entry skills:** What skills are students coming to the instruction with, both technical (e.g., basic computer, coding, computational) and psycho-social (collaboration, presenting, asking questions)?
- **Prior topic knowledge:** Specific to the content of the instruction, what do they already know coming in?
- **Attitudes toward content:** Are they excited to learn, or nervous about the content? How confident are they about their success in the class?
- **Attitudes about the delivery format:** Do they have attitudes, positive or negative, about the delivery format (e.g., lecture, flipped classroom, lab) and/or technologies (e.g., learning management systems)
- **Learning preferences:** Are learners comfortable in active learning experiences? (e.g., working with teams, interacting with the instructor, using technology while learning, etc.)
- **Group characteristics:** Looking at these considerations, how diverse is the group? Have they been together for instruction before? Do you have a small class or a very large class?

A simple online survey ahead of class can be very helpful in gaining an understanding of your audience.

3.2.3 Cultivate Student Study Skills and Learning Strategies

Do not assume that your students will have acquired the necessary study skills to effectively and efficiently learn from your course. From note-taking skills (i.e., finding patterns, discerning what is worth writing down, linking to readings, etc.) to multitasking (i.e., knowing when it is ok to keep their email open while studying and when they really have to focus) are skills that often have to be developed during learning experiences.

For most students learning through hands on tools, such as notebooks, will be quite different than their previous learning experiences—especially coming out of traditional high school or large-lecture undergraduate course. Interacting actively with technologies for the express purpose of learning (i.e., not just socializing with friends) is valuable—and yet not common.

Take time to work with your students to help them build the foundations for how they can most benefit from the experiences you are creating. Talk with them early in the instruction about how notebooks will be used and how they will have to adapt their study strategies (where, when, and how they study) to best learn the content and achieve the goals of the instruction.

3.2.4   Develop Your Content Strategy

The content of the instruction should support the instructional goal(s) you have identified for the course. Sometimes the goal(s) and the content are synonymous, but other times they are not. For example, if your goal is for students to be able to calculate conjoint probabilities, then your content will be synonymous and be on calculating conjoint probabilities. Whereas in other contexts, your goal may be to develop critical thinking skills in relation to logical fallacies, and the content you use to achieve this is the analysis of political discourse, which is not necessarily synonymous.

In both cases, the content guides why, when, and how notebooks can be used to achieve the goals. Remember that notebooks are solutions we can use to achieve our goals; they should not be confused with the goals themselves, even then they are closely related.

The content of your instruction is not limited to what you want to teach; it involves activities, exercises, feedback opportunities, and assessments. Each of these supports your goal(s) for instruction and benefits from the use of different tools within and outside of Jupyter notebooks. For example, if your content goals suggest that students require some specific knowledge (such as, what is a logic fallacy) before they are prepared to move on to another content topic, then you should assess that knowledge before you select which tools (e.g., a fill-in-the-blank item in a notebook, or a verbal question to the students in a classroom, or a self-reflection) is most appropriate for that step in your instruction.

Having a comprehensive outline of the goals and related content for the instruction is key to making good decisions about how, when, where, and why to use notebooks in your instruction.

3.2.5   Consider the Context of the Learning Environment

Instruction happens in many interesting contexts—but learning is not restricted to the instructional texts that we construct. Students learn before, during, and after instruction, and notebooks may be a integrated components for students in all those contexts (see below).

Within our instructional contexts, those times when we likely have the most influence on student learning, we also have lots of options for how we delivery learning experiences. As an instructor you can, for example, assign notebooks as pre-work for a flipped classroom approaches, and/or you can use the notebooks during an in-person or online class to demo or offer student practice, and/or notebooks can be used for homework, assessments, and resources that students can use later long after your instruction. All of these create different contexts for your use of notebooks.

Other aspects of context that you should consider when determining your use of notebooks including both those of the instructional environment and then the later performance environment of the learners.

3.2.5.1   Instructional Environment

Today the instructional environment can be quite complex and involve multiple aspects. For instance, the instruction may include online videos, short lecture in a classroom, and then group activities in a lab. Each of these are unique environments and their characteristics may influence how, when, where, and why you use notebooks.

- **Classroom:** How will students be engaged with the content? Given the layout of the room (e.g., small tables, lecture hall) are there opportunities for peer engagement?

- **Lab:** How does the physical structure of the lab environment offer opportunities for peer learning and sharing? What are the hardware and software tools available to students?
- **Flipped classroom:** What are required knowledge and skills that students must gain from pre-activities? What is the role of video in the flipped classroom? Are students prepared with the study skills and strategies for learning independently in the flipped classroom approach? Will all students have access to the materials (potentially in a notebook) prior to the class? Do students have access to computers outside of the classroom? How will pre-reading/pre-watching be incentivized and assessed in the course?
- **Online classroom:** Are students prepared with the study skills and strategies for learning independently in the online classroom?

3.2.5.2 Performance Environment

After the instruction your students will, hopefully, apply what you have taught them, and the environment in which they apply your instruction can vary greatly. You will want to consider how any differences between the instructional environment and the performance environment might impact on the ability of students to apply what they have learned.

- **Organizational/Managerial Support:** Will students be supported in their use of the instruction and tools that you are using?
- **Social Support:** If they learn through group projects, will they also be able to apply the instruction when working on their own?
- **Physical Aspects:** Will they have access to the same, or similar, tools and resources they have access to during the instruction? For example, if students only learn in notebooks, will they also be able apply the learning in an IDE that is used by their employer?

3.2.6 Crafting an experience: choose the right tool and approach for the task

As an instructor, you create many types of learning objects for your classroom. Jupyter notebooks can be used to present many types of information, from slideshows to book chapters to homework. You can imagine the large differences you might have with the information you present within a slideshow, a book chapter, a homework assignment, and a worksheet for an in class activity. Jupyter notebooks can be used for all these activities, yet this is all within the same type of document platform.

This freedom of expression available within the Jupyter notebook environment can add a lot of pressure and decision fatigue to the design process. A blank notebook you intend to make a lecture document within and a blank notebook for an in class activity each look the same when first created.

Content and design decisions should be driven by the purpose of the lesson and the needs of the audience. How many times have you asked someone a question and had them answer it in a completely unhelpful way? Perhaps you asked about how to implement something but they answered it conceptually, or you asked for reasoning about why something existed but they gave you a syntax information.

Learning goals and pedagogic practice are complex and often domain and skill specific, but decisions about them come back to the essential question: what are you trying to do here? As an instructor, you not only need to select the right technology for the learning experience, but the right activities to fulfill the learning mission of the day and the course.

Students are given a variety of experiences to aid in the learning process. Some topics are conceptual that students need time to experiment with and build intuition over; others are purely syntax that requires reference and practice to build their internalize knowledge over usage, and some fall in between.

The rhythm of the learning process sees students working independently, working together, reading, producing, listening, problem solving, and struggling. Each activity should serve a purpose as part of a larger

experience. The same way that restaurant or business owners craft an experience for their customer, you are crafting an experience for your students.

This design of this experience is not a simple process, and something driven by your expertise of your audience, domain, and your own ability to instruct. Through it all, your core perspective should be to incorporate elements that harmonize together to make an experience leading to your desired learning goal.

Perhaps for the section above add a block diagram or word cloud visual of the major decision considerations.

3.2.7   Transitioning to and from Jupyter

Although this book focuses on using Jupyter notebooks in education, we recognize that students have different backgrounds and familiarity with coding. Not all students are ready to jump right into notebooks from a traditional lecture-based classroom; while, other students may have significant programming experience and may be closer to being ready to transition to IDEs.

Based on your analysis of the goals and context you will want consider the appropriate instruction points for introducing and exiting from notebooks as a tool for achieving your goals. As your students gain experience with programming, they may be more interested in using an IDE. Likewise, you may choose to introduce a more traditional IDE first and introduce Jupyter notebooks later. Exercises that can support this transition include assignments where they do the same task in 3 or 4 different IDEs and then reflect on those experiences as they make decisions about which environments best supports their ambitions at the time.

3.2.8   Conclusion about teaching practices

In this chapter we have attempted to place the use of Jupyter notebooks into the process of good course design. Notebooks are a tool that many teachers can use to increase student engagement, participation, understanding, and performance—but that is not to say that every course and every lesson should use notebooks. Let the goals of your course define when and how to best use notebooks to achieve those goals.

Chapter 4

A catalogue of pedagogical patterns

4.1 Introduction

In this chapter, we present a collection of patterns that are particularly aligned with teaching and learning with Jupyter. Each pattern is targeted at specific learning goals, audiences, and teaching formats. With those in mind we describe each pattern and its pedagogical features that support the learning goals, present a practical example, and close each with any potential pitfalls you would want to be aware of.

4.2 Shift-Enter for the win

Description: Instead of reading a static chapter about a topic, the learners read and execute code, as well as potentially interact with a widget to explore concepts. Starting from a complete notebook, the instructor or learner runs through the notebook cell-by-cell by typing SHIFT + ENTER.

Example: The notebook (or a collection of notebooks) can be used as an alternative to a static textbook on a topic.

Learning goals: This pattern can be used to introduce a topic or promote awareness about a set of tools. Additionally, it can serve as documentation that provides a tour of an application programming interface (API).

Audience(s): Depending on the style of the notebook, this pattern can be used for a spectrum of programming abilities.

Format (lecture / lab / ...): This pattern can be used as an alternative to a static textbook. In a tutorial, a complete notebook can be used to provide a tour of an application programming interface (API) of a software package.

Features: One benefit of this approach is that learners have a complete working example which they can adapt or build from. It provides opportunity for richer interaction than a static textbook.

Pitfalls: This style does not prompt much engagement with students. Having a class that interactively works through a notebook can lead to some students finishing much faster than others (e.g., racing through SHIFT + ENTER). Breaking long notebooks into many smaller ones can help with the pacing in a lecture. Having a master notebook serve as the table of contents can then help students navigate through the class. Notebooks can be linked in a markdown cell as:

```
[Notebook 1](part1.ipynb)
```

4.3   Notebook as an app

Description: Notebooks can be used to rapidly generate user interfaces where students and instructors can interact with code through sliders, entry boxes, and toggle buttons. The code can run numerical simulations or perform simple computations, and the output is often a graph or image.

Example: In geophysics, a direct current resistivity survey involves connecting two electrodes to the ground through which current is induced. Current flows through the earth and the behavior depends upon the electrical resistivity of the subsurface structures; current flows around resistors and is channeled into conductors. At interfaces between conductors and resistors, charges build up, and these charges generate electric potentials which we measure at the surface of the earth. Each of these steps can be demonstrated through a simulation where students or the instructor builds a model, and views the currents, charges and electric potentials.

Related patterns: Top-down sequence

Learning goals: This approach can be effective for focusing on domain-specific knowledge and facilitating the exploration of models or computations.

Audience(s): This style can be effective for students with minimal programming experience as they do not have to read, write, or see the code.

Format (lecture / lab / ...): In lecture, this style of notebook can be used by an instructor to methodically walk through a concept step-by-step. It is also useful for promoting in-class engagement as students can suggest different parameter choices and instructors can adapt the input parameters based on students' questions.

In a lab or assignment, the notebook can be used as a "app" around which questions and exercises are built.

Features: Notebooks as apps can be used to promote engagement with students in lecture. In labs, assignments or in-class activities, this approach lowers the barrier-to-entry for students to explore complex models.

Pitfalls: It is important to have well-structured exercises and questions for students to address with the app. As with any app, simply asking students to play with it does not promote productive engagement.

In structuring an exercise for students, we recommend putting instructions and questions in a separate document rather than in the notebook. If students view the notebook as an app, they often want to interact with it rather than read it. By having instructions and questions that go alongside the notebook, they can have the app in view while reading.

This approach is not intended to be used for developing students' programming skills.

Enabling technologies: Widgets, domain-specific libraries such as simulation tools.

4.4   Win-day-one

Description: A win-day-one exercise brings learners to the answer quickly and concisely, almost like a magic trick, and then breaks down and methodically works through each of the steps, revealing the magician's tricks. It generally involves multiple notebooks: the first notebook being the "win" which shows the workflow end-to-end, and subsequent notebooks breaking down the details of each component of the workflow.

Example: To solve a numerical simulation using a finite volume approach, a mesh must be designed, differential operators formed, boundary conditions set, a right hand-side generated and then the system solved. Naturally, there are important considerations for each step. For even a moderately sized problem, sparse matrices are necessary in order to keep memory usage contained, the mesh must be appropriately designed in order to satisfy boundary conditions, and the solver needs to be compatible with the structure

of the system matrix. These details are critical for assembling a numerical computation, but if introduced upfront, they can overwhelm the conversation.

In a win-day-one approach, learners are first shown a concise example, in which many of the details are abstracted away in functions or objects. For example, methods such as `get_mesh`, `get_pde`, and `solve` abstract away the details of mesh design, creating differential operators and solving the set of equations. In subsequent notebooks, the workflow is tackled methodically, and the inner workings of each component discussed.

Related patterns: Top-down sequence, Proof by example, disproof by counterexample

Learning goals: This can be an effective approach for introducing complex processes, providing context for how each of the components fits together, and focusing attention.

Audience(s): This style can be effective for a spectrum of student audiences from those with some programming experience to those with significant experience.

Format (lecture / lab / ...): This can be effective for tutorials and workshops, and can be used over multiple lectures. This can be useful when introducing new topics to help hook students because they accomplish significant results early in the learning process.

Pitfalls: One hazard of the "win-day-one" is that the "win" is overwhelming (too much detail) or too magical (too little detail). An appropriate level of detail needs to be selected so that each of the components of the workflow is demonstrated, but at a high-level.

4.5 Fill in the blanks

Description: To focus attention on one aspect of a workflow, the scaffolding and majority of the workflow can be laid out and some elements removed with the intent that students (or the instructor during a demo) fill in those pieces. The exercise might be accompanied by a small test that the code should pass, or a plot, or value which the code should generate if correct.

Example: A fundamental concept in computing is the use of a for loop to accumulate a result. A fill-in-the-blank exercise demonstrating an accumulator could lay out the initialization, provide the skeleton of the for loop and include plotting code, with the aim being that students write the update step inside the for loop.

Related patterns: This pattern is similar to [Target practice] **TODO: missing section?**; a difference is that Target Practice often focuses on a bigger step in a multi-step process. Fill in the Blank exercises tend to be smaller and more immediate.

Learning goals: This pattern focuses attention on a component of a task and provides the benefit of demonstrating how that component fits into the bigger picture or a larger workflow. It can be an effective approach for taking students on a tour of an API, requiring that they use the documentation of the software, or for focusing attention on one aspect of an multi-step computational model.

Audience(s): This approach can be used with a range of students, from those who are first being introduced to computing concepts to those who have significant experience.

Format (lecture / lab / ...): Assignments and labs can adopt this approach (nbgrader can be used to help with marking). It can also be used in a lecture or tutorial setting where the instructor demo's how to fill in the blank.

Pitfalls: Some students don't find this approach engaging. In particular, if the exercise is too simple for the level of programming competency of the students, it can be perceived a "make-work" task.

4.6   Top-down sequence

Description: Particularly in STEM, the default sequence of presentation is bottom-up, meaning that we teach students how things work (and sometimes prove that they work), before students learn how to use them, or what they are for.

Notebooks afford the opportunity to present topics top-down; that is, students learn what a tool is for and how to use it, before they learn how it works.

Example: In digital signal processing, one of the most important ideas is the discrete Fourier transform, which depends on complex arithmetic; in a bottom-up approach, we would have to start by teaching or reviewing complex numbers, which is not particularly engaging.

In contrast to writing the mathematics on paper, in a notebook students can use a library that does the discrete Fourier transform for them, so they understand what it is used for, and see the value of learning about it, before we ask them to do the work of understanding it.

Learning goals: This pattern is useful for shifting students' focus from implementation details to domain content.

Audience(s): This pattern can be effective with students who have limited programming skills, as they can use a library and see the results without writing much, if any, code.

Format (lecture / lab / ...): This pattern can be used in a single class session or homework, or spread out over the duration of a course; for example, students could use a tool on the first day and find out how it works on the last.

Features: Shows students value and rewards their attention quickly (see "Win Day One").

Pitfalls: A potential hazard of this pattern is that students might be less motivated to learn how the tool works if they think they have already understood what it is for and how to use it. This hazard can be mitigated by making obvious the additional benefit of understanding how it works (assuming that there actually is one—it is not enough to assert that knowing how it works is necessarily better).

4.7   The API *is* the lesson

Description: When students work with a software library, they are exposed to functions and objects that make up an application programming interface (API). Learning an API can be cognitive overhead; that is, material students have to learn to get work done computationally, but which does not contribute to their understanding of the subject matter. But the API can also be the lesson; that is, by learning the API, students are implicitly learning the intended content.

Example: In digital signal processing, one of the most important ideas is the relationship between two representations of a signal: as a wave in the time domain and as a spectrum in the frequency domain. Suppose the API provides two objects, called Wave and Spectrum, and two functions, one that takes a Wave and returns a Spectrum, and another that takes a Spectrum and returns a Wave. By using this API, students implicitly learn that a Wave and a Spectrum are equivalent representations of the same information; given either one, you can compute the other.

Related patterns: Top-down sequence

Learning goals: This pattern is useful for shifting students' focus from implementation details to domain content.

Audience(s): This pattern is most effective if students have some experience using libraries and exploring APIs.

Format (lecture / lab / ...): This pattern TODO: complete

Pitfalls: A hazard of this pattern is that students sometimes perceive the costs of learning the API and do not perceive the benefits. It might be necessary to help them see that learning the API is part of the lesson and not just overhead.

4.8 The world is your dataset

Description: Notebooks provide several ways to connect students with the world beyond the classroom: one simple way is to collect data from external sources. Data is available in many different formats that require different software tools to collect and parse.

If a dataset is available in a standard format, like CSV, it can be downloaded from inside the notebook, which demonstrates a good practice for data integrity (going to the source rather than working with a copy) and demystifies the source of the data.

For data in tabular form on a web page, it is often possible to use Pandas to parse the HTML and generate a DataFrame. Also, for less structured sources, tools like Scrapy can be used to extract data, "scrape", from sources that would be hard to collect manually, and to automate cleaning and validation steps.

Examples: Datasets like the National Survey of Family Growth are available in files that can be downloaded directly from their website, but the terms of use forbid redistributing the data. So the best way for an instructor to share this data is to provide students with code to input into a notebook cell, which, when executed, will download the data set the first time the student runs the notebook.

Many Wikipedia pages contain data in HTML tables; most of them can be imported easily into a notebook using Pandas.

Sources of sports-related statistics are often embedded in large networks of linked web pages. Tools like Scrapy can navigate these networks to collect data in forms more amenable to automated analysis.

Audience(s): Students with limited programming experience can work with datasets in standard formats, but scraping data might require more programming experience.

Format (lecture / lab / ...): This pattern lends itself to more open-ended project work where students are responsible for identifying data sources, collecting data, cleaning, and validating, but it can be adapted to more scaffolded work (see Target Practice).

Features: Contributes to students' feelings of autonomy and connectedness.

Pitfalls: A hazard of this pattern is that students can spend too much time looking for data that is not available. They might need coaching about how to make do with the data they can get, even if it is not ideal.

Enabling technologies: Pandas, Scrapy, R, ROpenSci packages

4.9 Connect to external audiences

Description: This is in some sense the opposite of "the world is your dataset." Here the goal is to take a workflow or computational exploration and share it with the world so others can see it, learn from it, reuse and remix it.

Examples: Your students are doing an observational astronomy lab where they take data from a telescope of a transiting exoplanet (a planet around a star other than our Sun) and they examine the lightcurve to learn about the planet. The students present the lab as a Jupyter notebook with a reproducible workflow that starts with reading in their data (images), walks through cleaning and reducing the images, and then performs photometry on the host star to produce a lightcurve. The end product is a plot showing the star's

brightness, dimming just slightly as the unseen planet comes between the star and earth. Elated with their result, the students want to share their data and workflow so anyone else can redo the analysis.

Learning goals: Reproducibility is an important part of the scientific process. Having completed the primary scientific analysis that was the goal of the lab (obtaining a lightcurve of a transiting exoplanet), the students now can learn reproducible science practices by hosting their notebook on a webserver (e.g., Github) along with the data. An essential part of making the notebook reproducible will be ensuring that the notebook clearly lists the needed dependencies.

Audience(s): All students—everyone should learn about reproducibility.

Format: A self-contained notebook hosted on a webserver.

Features: Teaches students about reproducible science workflows.

Pitfalls: You need to be clear about the library requirements needed to run the notebook. Also, since the data files are likely separate from the notebook, it is possible for copies of the notebook to get shared without the data. Students may also be shy or fearful of showing their work publicly, so explaining the benefits may be needed to curtail their worries.

4.10   Twiddle, tweak, and frob

Description: Students are given a notebook with a working example. They start by reading the text, running the code, and interpreting the results. Then they are asked to make a series of changes and run the code again; the changes can be small (tweaks), medium-sized (twiddles), or more substantial (frobs). For the origin of these terms, see *The New Hacker's Dictionary* (Raymond, 1996).

Offering manipulations on a range of scales allows students to interact with notebooks in ways that suit their background and styles. Students who feel overwhelmed by the technology can get started with small, safe changes, enjoy immediate success, and work their way up. Students with more experience or less patience can make more radical changes and learn by "destructive testing".

Examples: In machine learning there are many steps to implementing an effective algorithm:

- Understand a problem
- Identify the proper machine learning algorithm to create the desired results
- Identify data and feature sets
- Optimize the configuration of the machine learning algorithms

We can write machine learning notebooks that allow students to modify parameters and interact at multiple levels of detail:

- **twiddle** hyper parameters to quickly see minor improvements in the results.
- **tweak** feature sets to create new models for a bigger impact
- **frob** by replacing the machine learning algorithm with a new algorithm or a new version

This pattern is particularly suitable for examples with a lot of parameters.

Learning goals: This pattern helps students acquire domain knowledge by seeing the relationship between parameters and the effect they have on the results. It can also help students learn new notebook use patterns

This pattern is similar to "The notebook is an app"; a difference is that in this pattern the code is more visible to the students, which can help orient them if they will make bigger changes in the future.

Audience(s): This pattern can work with students who have no programming experience; they only need to be able to edit the contents of a cell and run a notebook. It can also work with students who have no background with the domain; they can learn about the domain by exploring the effect of parameters.

Format (lecture / lab / ...): This pattern lends itself to a workshop format, where students are guided through a notebook with time-boxed opportunities to experiment. It also lends itself to pair programming, where a navigator can suggest changes and a driver can implement them.

Features: Can help students overcome anxiety about breaking code, and build comfort with self-directed exploration.

Pitfalls: One hazard of this pattern is that students might have trouble getting started, so you might suggest a few examples; however, another hazard is that, if you provide examples, students will do what they are told and fail to explore. A third hazard is that the changes a student makes might be too unorganized; in that case the effect of the parameters might be lost in the chaos.

Enabling technologies: Ideally the students should work in some kind of version control system that lets them revert to a previous version if they break something (and don't know how to fix it). Note that undo and redo, Ctrl-z and Ctrl-y, can be used to traverse deep history within each cell, but not across cells.

4.11   Coding as translation

Description: Converting mathematics to code is a critical skill today that many students, especially those without strong programming backgrounds, struggle to do. Explicitly taking an equation and translating it step-by-step to the code can help these students make the transition to attaining this skill.

Example: Say you wanted to show the translation of matrix-vector multiplication from *equation* to a numerical computation. This would involve setting up and explaining the mathematics and suggesting replacing the sums with loops and initializing the sum properly.

Learning goals: Translating mathematics to code (and vice versa)

Audience(s): Learners who understand the theory but struggle with the programming side of things.

Format: This type of pattern is often best served as a notebook with some explanatory text and possibly some scaffolded code so that a student can focus on the critical areas. This can be done as easily as a lab exercise or in lecture with perhaps some time held out for the students to solve it themselves before moving forward. It is critical to this pattern that there is a clear connection between the mathematical symbols (such as the summation) to the code (such as the for loops).

Features: This pattern can work to lower the barrier for students with low programming knowledge to take on more complex tasks.

Pitfalls: If the exercise is not properly scaffolded, namely it is too complex, students can be turned off. This is especially true if the code example is too complex with too many steps. For instance avoiding compound operators in the example above (`+=`) can help student retention.

4.12   Symbolic math over pencil + paper

Description: Your objective is to convey an understanding of a physical system governed by a complicated mathematical system. Working out the algebra is necessary to uncover the fundamental behavior of the system, but how to do the algebra is not the goal of the lesson. In this case, you want to see the algebraic result and then teach the students the underlying meaning of the system,

Example: The Euler equations for hydrodynamics are a system of partial differential equations governing conservation of mass, momentum, and energy. Their mathematical character admits wave solutions, and the eigenvalues and eigenvectors of the system in matrix form are important to understanding the physical behavior of the density, pressure, velocity, etc., in the system. Working out the eigenvalues with pencil and paper is tedious, and not the objective of the lesson. In this case, we can use a symbolic math library like

SymPy (Meurer et al., 2017) to do the mathematical analysis for us, finding the eigenvalues and eigenvectors of the system, and we can then use this result to continue our theoretical discussion of the system.

Learning goals: Students will see how to do symbolic math that arises in their theoretical analysis.

Audience(s): STEM students that want to focus on understanding a mathematical system without worrying about the algebraic details.

Format: This works well as a notebook that acts as a supplement to the main lecture. Since the goal of the lecture is the theory, the notebook can [TODO: complete]

Features: Abstracts the details of mathematics that is secondary to the discussion at hand into a separate unit that students can explore on their own.

Pitfalls: This only works well in the case that the algebra is not essential to the main learning goals, but rather is simply something that must be done to get to the main goal.

4.13   Bug hunt

Description: The instructor provides a notebook with code that contains deliberate bugs. The students are asked to find and fix the bugs. Automated tests might be provided to help students know whether some bugs remain unfixed.

Example: TODO

Learning goals: This pattern helps students develop programming skills, especially debugging (of course); it also gives the practice reading other people's code, which can be an opportunity to demonstrate good practice, or warn against bad practice. It can also be used to teach students how to use debugging tools.

Audience(s): This pattern requires students to have some programming experience.

Format (lecture / lab / ...): This pattern can be used for in-class activities or homework.

Features: Can be engaging and fun; develops important meta-skills.

Pitfalls: The bugs need to be calibrated to the ability of the students: if they are too easy, they are not engaging; if they are too hard, they are likely to be frustrating.

4.14   Proof by example, disproof by counterexample

Description: In many classes, students see general results derived or proved, and then use those results in programs. Notebooks can help students understand how these results work in practice, when they apply, and how they fail when they do not.
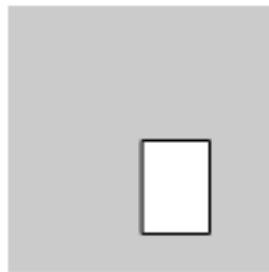
Example: In statistics, the Central Limit Theorem (CLT) gives the conditions when the sum of random variables converge to a Gaussian distribution. Students can generate random variables from a variety of distributions and test whether the sums converge and how quickly.

The classical Gram–Schmidt is unstable while the modified method is stable. Students can find matrices for which this instability produces obviously unusable results. They can also find matrices for which modified Gram-Schmidt produces unusable results due to its lack of backward stability, and this can be used to motivate Householder factorization and discussion of backward stability.

Some numerical methods for PDE converge with an assumption on smoothness of coefficients. Students can show how violating these assumptions leads to erroneous solutions, thus motivating discussion of conservative/compatible methods that can converge in such circumstances.

Figure 4.1: Figure: a first sketch using Calysto Processing, a Java-based language designed for creating art.

Learning goals: This pattern is primarily useful for developing mathematical or domain knowledge, but students might also develop programming experience by writing code to run the examples and test the outcomes. This is especially true if the space of (counter-)examples is "small", such that principled exploration (e.g., by finding an eigenvector, running an optimization algorithm, or searching a dictionary) is beneficial.

Audience(s): This pattern requires students to have some programming experience.

Format (lecture / lab / ...): This pattern can be used for in-class activities or homework.

Features: Helps students translate from theoretical results to practical implications, and to remember the assumptions and limitations of theory.

Pitfalls: This pattern requires additional time and student effort on a topic that might not deserve the additional resources.

4.15 Hello, world!

Description: In some situations (such as the first day of class of a very introductory course) you may wish to do no more (and no less) than build confidence in the students' abilities to be able to write a first computer program. Traditionally, the first program written was a "hello, world" program: a program that did nothing but display the text "hello, world" on the screen. However, these days students can have much more fun, and step fully into the creative world of computing with very little instruction.

Example: Draw a rectangle. Change the numbers, run it again, and see what happens.

Learning goals: Reduce stress, build confidence, connect onto their personal lives. Requires that they do learn the basics of Jupyter including: log in, open a new notebook, enter the provided code, and execute it. Often leads to a very animated, active learning classroom activity ("How can I change the colors?", "How do I draw a circle?", etc.)

Audience(s): Beginning students.

Format (lecture / lab / ...): First day of class, in-class exercise. Build on what students already know from typing and reading (e.g., cut and paste, read top-to-bottom).

Features: Open ended, creative, fun.

Pitfalls: Works best when used with a pre-installed Jupyter (see the relevant chapter). Rather than telling students that they can do it, just do it. As a first assignment, to cut down on the vast possibilities, we suggest limiting the palette of options. For example, restrict their drawings to use only a single shape, such as rectangle or triangle. We suggest having the students draw something in their life that is important or meaningful to them. We suggest discussing the coordinate grid for the first assignment and sketching an idea on paper first.

4.16   Test driven development

Description: The instructor provides tests written in a unit testing framework like `unittest` or `doctest`; students write code to make the tests pass.

Example: TODO: Necessary?

Learning goals: Helps students learn a good software development process.

Audience(s): This pattern requires students to have some programming experience.

Format (lecture / lab / ...): This pattern can be used for in-class activities or homework.

Features: Helps students focus on the task at hand and know when they are done (at least to the degree that the tests are complete).

Pitfalls: Some Python unit testing frameworks are not designed to work with notebooks, and can be awkward to use. On the other hand, `nbgrader` [TODO: add cross reference to nbgrader] supports automated testing of the code students write in notebooks; in that environment, the tests are not visible to students, which may or may not be a bug.

This pattern required the overhead of teaching students about the unit testing framework. Students working to make tests pass can lose their view of the big picture, and feel like they have been robbed of autonomy. This type of exercise is best used sparingly.

4.17   Now you try (with different data or process)

Description: Students start with a complete working example provided by an instructor and then they change the dataset or process to apply the notebook to an area of their own choosing. This method can allow more or less fluctuation depending on the skills of the students. For example we can allow students to select new datasets from a list that ensures the cells of the notebook will all still work or we can give them freedom to try new data structures or add new processes to break the notebooks and learn as they go through the process of fixing the broken cells.

Example: An instructor designs a lesson in exploratory data analysis to scrape the critics reviews for a specific movie from a particular movie review website and then provide some simple visualizations. The students have a few options:

1. Green Circle - replace the movie name and pick any movie they want and then step through the new notebook and see the new results.
2. Blue Square - adjust the notebook to scrape user reviews rather than critics reviews and then fix any data parsing problems.
3. Black Diamond - add different visualizations tailored to explore the user reviews (as opposed to the initial visualizations that are tailored for the critic reviews).

There are various ways to test the properties of numerical methods. For example, students can use the method of manufactured solutions to test the order of accuracy for a differential equation solver. They can also measure cost as the resolution is increased and present the results in a way that would help an analyst decide which method to use given external requirements (e.g., using accuracy versus cost tradeoff curves).

Related patterns: Top-down sequence

Learning goals: This pattern allows students to apply their knowledge

Audience(s): This pattern can be tailored for students with more or less experience even in the same course.

Format (lecture / lab / ...): This pattern is best in a lab or an interactive tutorial

Pitfalls: A hazard of this pattern is that students may go completely off the rails and chose datasets or new processes that have not been tested and will not work in the timeframe allowed.

4.18 Two bites at every apple

Description: This pattern involves writing an activity that can address multiple audiences from different perspectives at the same time. This can be powerful when addressing a mixed audience of students.

Example: Say you have a group of students, some of whom are computer science students and some of whom are physics students, and ask them to come up with two expressions for computing the centroid of an area. The computer science students will be tasked with a description that involves adding up discrete pieces of areas with for loops and the physics students tasked with using the integral definition. When the students come up with their expressions they can then pair up with someone from the other background where they can attempt to explain how their approach matches the other and compare their final answers.

Learning goals: Ability to translate from one field/language to another. Explain complex topics to someone from a different field.

Audience(s): Groups which are composed of disparate backgrounds.

Format: This format involves both individual and group work but can be used in a lab or lecture setting. The basic notebook would include an overview of the problem and then pose questions whose answer is the same but is worded for the different audiences. There can be a single notebook that contains both questions so that students can fill in their peers solution once they understand it or there can be separate notebooks for each group so they do not get distracted by the other question.

Features: Group work and peer teaching has been shown to be effective at not only reinforcing student knowledge but also at introducing students to new concepts.

Pitfalls: It can be difficult to construct questions for each audience that require equal amounts of difficulty.

4.19 Replace analysis with numerical methods

Description: Some ideas that are hard to understand with mathematical analysis are easy to understand with computer simulation and numerical methods.

In the usual presentation, students see and learn to do mathematical analysis on a series of simple examples, and resort to numerical methods only when necessary. In an alternative pattern, students skip the analysis and start with simulation and numerical methods, optionally visiting analysis after gaining practical experience.

Examples: In statistics, hypothesis testing is a central idea that is notoriously difficult for students to understand. Students learn methods for computing p-values in a series of special cases, but many of them never understand the framework, or what a p-value means. The alternative is to compute sampling distributions and p-values by simulation; anecdotally, many students report that this approach makes the framework much clearer. Such simulation can also be used to drive home points about misconceptions held by most students and instructors (Haller & Krauss, 2002).

Similarly, in queueing theory, there are a few analytic results that apply under narrow conditions; when those conditions don't apply, there are no analytic solutions. However, queueing systems lend themselves to simulation and visualization, and in simulation it is easy to explore a wide range of conditions.

And again, with differential equations, there are only a few special cases that have analytic solutions; the large majority of interesting, realistic problems don't.

Learning goals: This pattern is primarily about helping students see the big ideas of the domain more clearly, but it is also a chance to develop their programming skills. It also provides students with tools that are likely to be needed if they encounter similar problems in the real world, where analytic methods are often inapplicable, fragile, or complicated to use effectively.

Audience(s): This pattern requires students to have some comfort with programming, although it would be possible for them to get some of the benefit from seeing examples without implementing them. Non-programmers can use this pattern via prepared notebooks; see Win Day One.

Format (lecture / lab / ...): This pattern can be used for in-class activities or homework.

Features: Students can understand general ideas without getting bogged down in the details of special cases; and they are able to explore more interesting and realistic examples.

Pitfalls: If students are not comfortable programmers, they can get bogged down in implementation details and debugging problems, and miss the domain content entirely. It is important to scope the implementation effort to suit the full range of students in the class. Pair programming can help mitigate these problems, especially if every pair has at least one student with programming skills, and if students are coached to pair program effectively (without letting the more experienced student dominate).

4.20   Code reviews

Code reviews involve a student or instructor providing feedback on someone else's code. This pattern involves peer work as well as a means for providing feedback to students on topics other than correctness of their code but also on code readability and styling.

Example: Present a problem to students that they must write a solution to, say computing the square root of a number without using a built-in function but have them write a test for their function that uses a built-in function to compute the answer. After they are finished have the students pair up and perform peer reviews of each other's code, commenting not only on the way they solved the problem, such as making up a list of pros and cons of their approaches, but also on the readability of the code.

Learning goals: Learn to read and understand someone else's code. Learn to write readable code.

Audience(s): Any group of students who are involved in coding.

Format: Once a suitable problem is formulated in a notebook (or simply a script) then in-class review, as with the above example, can work for peer reviews. Alternatively students can upload their notebooks/scripts to a platform such as GitHub and the code reviews can be done using the tools available there. Sufficient scaffolding must be provided so that students understand the process, how to make constructive comments and why the process is important. If an instructor wants to review and provide feedback notebooks/scripts can be collected and commented on with a similar explanation to students as to how they are going to be graded (if they are).

Features: This pattern leads to not just feedback for the person who wrote the code but also for the reader. Code review is also a critical piece of the software development process used in industry providing students with a view of the process. This can also have the result of making sure that a student's code is readable via appropriate code styling, commenting and documentation.

Pitfalls: Students need to be properly informed as to how the code reviews will impact their grades, especially if peer review is used. Notebooks on GitHub are not as easily reviewed as scripts.

4.21 Adversarial programming

This pattern involves participants writing a solution to a problem and tests that attempt to make the written solution fail. This pattern can be done in many ways including having students complete the tasks and pair up and exchange solutions/tests or having the instructor writing the solution and the students then write the tests.

Example: Students are tasked to write a function that finds the roots of a polynomial specified via some appropriate input. They are also asked to write a set of tests that their function passes and fails on. When students have completed these tasks they then exchange their notebooks and use the tests they wrote on their peer's function. Finally they will discuss any differences in their approaches and whether they can come up with ways to not fail each other's tests or if the tests provided are invalid.

Learning goals: Learn to write unit tests. Think critically on how an adversary might break their solution.

Audience(s): Any group of students who are involved in coding.

Format: Decide on a sufficiently complex problem that may have non-trivial tests written for it and write up the question in a notebook. Then as an in-class activity or lab start the discussion regarding the tests. If appropriate the instructor can collect notable tests written by students and also share those.

Features: Provides a means for students to think critically about a problem they are solving and how someone might break their solution. Also can provide a learning activity with a form of competition involved, which can then lead to an award system if desired.

Pitfalls: With competition come dangers if students are not properly scaffolded so that they can provide constructive feedback. Some problems and/or solution strategies are vulnerable to many corner cases, leading to tedious whack-a-mole or fatalism that may distract from learning objectives.

4.22 There can be only one

Description: This pattern involves creating a competition between individual students or teams of students. Clear goals and metrics need to be defined and then students submit notebooks that are scored and evaluated. Competitions can span months or be completed in a single class.

The Jupyter ecosystem support for reproducibility and data sharing make it a great environment for creating healthy competitions. Kaggle is a site that hosts many machine learning competitions using Jupyter as its underlying infrastructure and is a great place for advanced students to extend their knowledge with a chance of winning cash prizes and solving current world problems.

Examples: Identify a machine learning problem and a labeled dataset your students can use to train their model. Then select an evaluation metric and detail your problem statement and rules. Finally, launch your competition and allow your students to submit their notebooks and post their results on a public leaderboard.

Learning goals: Creative problem solving is a key aspect of this pattern. In addition, if the competition is team-based then the students will learn how to work in groups and communicate effectively and share responsibilities.

Audience(s): Students can benefit from healthy competition and working in teams, however it is critical that a safe, fun, and engaging environment is created. Advanced students can be pointed toward kaggle or other public competitions that may be in their area of interest and give them a chance to test their skills in the real world.

Format: A competition can be defined with any metrics and rules and can be run in multiple ways. The students can help define the rules or a simple vote can decide the winners. For a more formal competition instructor can host free competitions for their class hosted by Kaggle. https://www.kaggle.com/about/inclass/overview

Features: Teaches students about creative problem solving and teamwork.

Pitfalls: Creating a fair competition is not trivial and considerations regarding data, metrics, rules, and scoring may be time consuming. Competition is risky business and effort should be made to ensure both winners and losers enjoy the experience.

Chapter 5

Jupyter ecosystem: notebooks

The Jupyter system supports over 100 programming languages (called "kernels" in the Jupyter ecosystem) including Python, Java, R, Julia, Matlab, Octave, Scheme, Processing, Scala, and many more. Out of the box Jupyter will only run the IPython kernel, but additional kernels may be installed. Additional language support continues to be added by the open source community and https://github.com/jupyter/jupyter/wiki/Jupyter-kernels is the best source for an up-to-date list. Note that these projects are developed and maintained by the open source community and exist in various levels of support. For example, some kernels may be supported by a vast number of active (and even paid) developers, while others may be a single person's pet project. When trying out a new kernel, we suggest exploring a kernel's community of users and documentation to see if it has an appropriate level of support for your (and your students') use.

Jupyter's kernel flexibility allows instructors to pick the right language for a particular context. For example instructors may use Python to teach programming, while switching to R to teach statistics, and then perhaps Scala to teach big data processing. Regardless of the language chosen, the Jupyter interface remains the same. Thus, some cognitive load can be lessened when using multiple languages in, or across courses (e.g., the user interface doesn't change at all between the student's Digital Humanities and Biology courses). Of course, if you can remain in one language in a course, that is often easier for the student.

5.1   Resources for authoring Jupyter notebooks

Before embarking on writing notebooks for your course, we recommend that you look around on the internet for related courses. There may be a similar course for which an instructor has already generated notebooks that you can use or adapt for your course. In many cases, instructors are very happy to have collaborate on open source educational resources and have the resources be used by other instructors. The following sections will be oriented towards Python simply because it is currently the language with the largest Jupyter feature support.

5.1.1   Accessing documentation in the notebook

One of the best features of quality libraries is their documentation, which students and other users will likely consult regularly. From a notebook cell, the TAB key autocompletes (or gives completion tips) and SHIFT-TAB brings up full documentation. Similarly, using a question-mark after a method or function will bring up the documentation after the cell is run, as shown in Figure XX.

Using this feature in class during live coding and modification or while just explaining how code works helps increase student comfort and enable them to work effectively with libraries.

```
Docstring:
print(value, ..., sep=' ', end='\n', file=sys.stdout, flush=False)

Prints the values to a stream, or to sys.stdout by default.
Optional keyword arguments:
file:  a file-like object (stream); defaults to the current sys.stdout.
sep:   string inserted between values, default a space.
end:   string appended after the last value, default a newline.
flush: whether to forcibly flush the stream.
Type:       builtin_function_or_method
```

Figure 5.1: A question mark can be used after a method or function and the cell run in order to bring up the documentation.

5.1.2   Widgets

Widgets provide the opportunity a for learners and instructors to interact with code outputs, such as charts and tables, interactively. Widgets are simple and quickly generated "mini" Graphical User Interfaces (GUI) that give the notebook user access to slide bars, toggle buttons, and text-boxes. They can be used in conjunction with code, allowing a change of mindset from programming as a primary goal to exploring a model or computation as the primary goal. Alternatively, the code can be hidden and the widgets used to create a notebook "app" that might connect input parameters with a simulation and a plot.

Currently only a small subset of kernels have widget functionality. For example, the reference implementation of widgets are the Jupyter-Python widgets (https://ipywidgets.rtfd.io). It includes widget components to generate and display sliders, progress bars, text boxes, check boxes, toggle buttons, etc. Many popular visualization tools, such as matplotlib, Plotly, leaflet.js, three.js, have Jupyter-Python widget implementations. The documentation contains an up-to-date list of all of the widgets and their variations. Additionally, the `interact` method wraps allows you to wrap a function, which might be a simple computation or a complex simulation that produces a plot, and provides widgets for each of the inputs to the function. Figure XX shows a simple example of a sinusoid plot whose frequency is controlled by a slide-bar. Another kernel that has some widget functionality is C++ (https://github.com/QUantStack/xwidgets).

In addition to the IPywidgets library, the ipyleaflet library (https://ipyleaflet.rtfd.io) enables an interactive map to be displayed in a notebook.

5.1.2.1   Example

```
from ipyleaflet import Map
Map(center=[34.6252978589571, -77.34580993652344], zoom=10)
```

For the ambitions reader, there are resources available for you to write your your own custom widgets. The widget cookie cutter project (https://ipywidgets.rtfd.io) is a good place to start.

5.1.3   Magics

Magics are meta-commands that only function within Jupyter and allow a user to access features that are language/kernel specific. For instance, the IPython kernel provides a number of magics that can be useful while developing Jupyter notebooks that uses Python as the primary language. These are documented and

Figure 5.2: Here, a simple slider is used to interactively change the variable k in our function as we plot it.

```
from ipyleaflet import Map

Map(center=[38.8997, -77.0486], zoom=17)
```
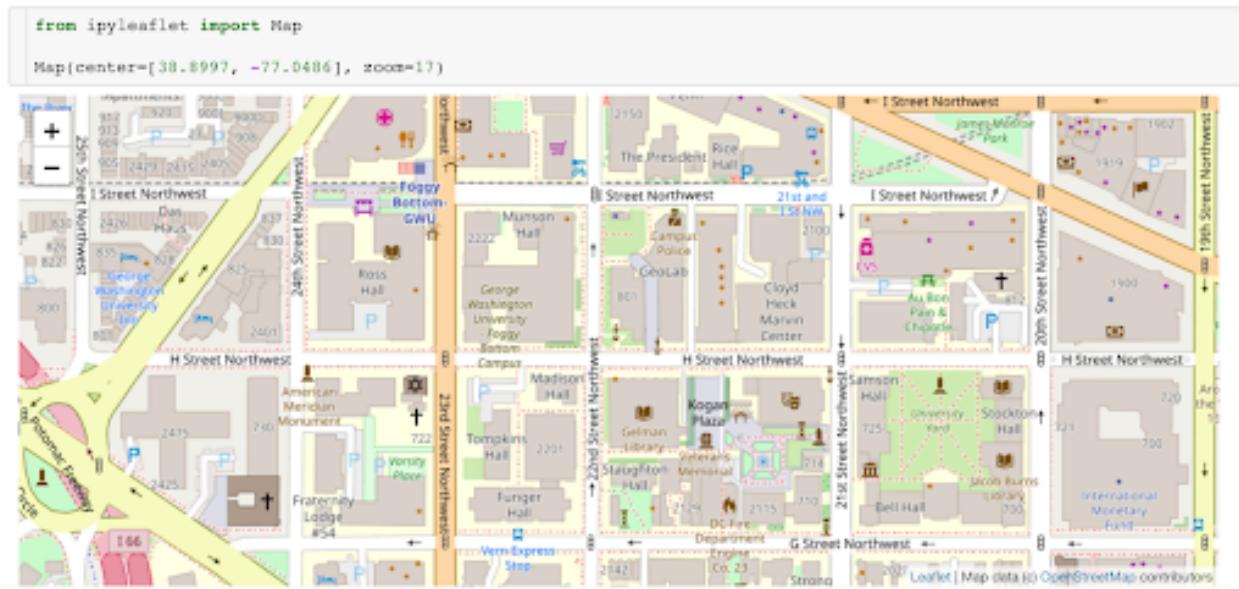
Figure 5.3: Interactive map widget with `ipyleafletalt_text`.

we will only call out a few of these here. There are many other magics available for different kernels but the downfall of these are that they are specific to Jupyter so does not translate to scripts and in a teaching environment we generally recommend to be sparing in their use so as to not obfuscate what is happening. Magics always begin with either a single `%` for single line commands or double `%%` for applying a command to an entire cell. Many magics can be used with single or double percents, but some do not.

- `%run` allows for running external scripts, capturing output and displaying it in the notebook, e.g. `%run my_script.py`.

- It is common to use matplotlib for visualization. In Jupyter, the magic `%matplotlib` allows the matplotlibs resulting figures to be displayed in the notebook. `%matplotlib inline` and causes static images to be embedded in the notebook and `%matplotlib  notebook` causes interactive images (zooming, panning, etc) images to be embedded.

- `%time` times the execution of the Python expression following the command, e.g.   `%time sum(range(1000))`.

- `%timeit` is is similar to `%time` with the difference that it runs the expression multiple times and reports the average execution time.

- `%reset` removes all user defined variables along with input and output. Magics often have "flags", following the Unix command pattern. For example, `%reset -s` is a soft reset and only removed user defined variables. These commands can be useful to help to overcome out-of-order execution problems.

A good example of a magic operating on the entire contents of cell is the `%%HTML` magic that will force the cell to be interpreted as HTML and render it appropriately. You can also use magics to call other languages while running the IPython kernel. For example, you can run R code from within an IPython notebook by using the `%%R` magic.

5.1.4   Notebooks under version control

Keeping notebooks under version control is a great way to not only keep track of changes to your created content but can also allow for sharable content. In a course where multiple people are contributing to the development of notebooks for the course, using version control in conjunction with a platform like GitHub,

allows authorship to be tracked and provides communication tools for reviewing new contributions or outlining requested development for a new assignment, activity, etc. Another advantage of using version control is that there are services that will provide rendered views of your notebooks that are publicly accessible. GitHub for instance will show a rendered version of the notebook rather than the direct text that a notebook is comprised of. There are a few pitfalls with this with LaTeX rendering as platforms do not always render the notebooks the same as they would appear in an active Jupyter interface.

There are a few caveats to keeping notebooks under version control that should be kept in mind. First of all the output will be stored in the repository, including images, if clear output is not used before committing changes. This can make detecting changes difficult as changes in output will be detected when nothing has actually changed content wise. The tracked notebooks also can become large if output is tracked. Even with clearing the output, reading through changes can be difficult due to the format of the notebook (notebooks are plain-text files and the file format is represented as JSON). This can be an issue if you are storing notebooks in a version control system, like git, and you wish to see differences between versions. For more information, see: https://github.com/jupyter/nbdime and https://github.com/mwouts/jupytext.

5.1.5 Testing notebooks

Before distributing notebooks, and in particular if you are working with multiple contributors to the course material, testing the notebooks before they are distributed to students or used in a live demo can help mitigate unexpected bugs. At a minimum, you can test that the notebook executes cleanly from top to bottom by restarting the kernel and running all cells from top to bottom. This can be done from the menu (Restart + Run all).

Though it requires a bit more setup, tests can be run automatically using a continuous integration service, such as TravisCI (https://travis-ci.org). This will require executing the entire notebook via the command line, for example `jupyter nbconvert --ExecutePreprocessor.enabled=True --to=html my_notebook.ipynb` will execute the notebook (same as pressing "Restart + Run All") and then convert it to HTML. These services can be connected to GitHub so that any time that the notebooks are changed, the tests are run automatically. Simplifying this process is an area that is under development in the open source community. The package https://github.com/opengeophysics/testipynb provides an easy way to test notebooks.

5.1.6 Essential Python libraries

The purpose of this section is to introduce some of the most widely used packages within the Python ecosystem. Python is not the only language that can be used in the notebook; over 90 different kernels enable different programming languages be used. We discuss the use of other languages in the section: Using other languages. Python has a large open-source community which develops and maintains an ecosystem of over 150,000 software packages, making it a common language of choice in many disciplines.

The core Python library (https://docs.python.org) contains data types such as lists and dictionaries, as well as core functionality such as arithmetic operators and simple file parsers. Most tasks can be achieved with core python, however, they are often made easier with higher-level libraries. These libraries are particularly useful for scientific computing with Python.

Among the vast number of packages in the Python ecosystem, Numpy, Scipy, Matplotlib and Pandas are among the most commonly used. Numpy (http://www.numpy.org/) is a fundamental library for numerical and scientific computing with Python. It contains data structures for numerical arrays, tools for linear algebra, and random number capabilities. Scipy (https://docs.scipy.org/) contains a variety of functionality for common scientific computing tasks, such as optimization, interpolation, statistics and signal processing. It also includes fundamental constants from many disciplines such as the speed of light as well as data structures for sparse matrices. Matplotlib (https://matplotlib.org/) is the core plotting library for python and can be used inline in the notebook with the `%matplotlib notebook` or `%matplotlib inline` cell magic.

▾ **TensorFlow execution**

Colaboratory allows you to execute TensorFlow code in your browser with a single click. The example below adds two matrices.

$$\begin{bmatrix} 1. & 1. & 1. \\ 1. & 1. & 1. \end{bmatrix} + \begin{bmatrix} 1. & 2. & 3. \\ 4. & 5. & 6. \end{bmatrix} = \begin{bmatrix} 2. & 3. & 4. \\ 5. & 6. & 7. \end{bmatrix}$$

```python
import tensorflow as tf

input1 = tf.ones((2, 3))
input2 = tf.reshape(tf.range(1, 7, dtype=tf.float32), (2, 3))
output = input1 + input2

with tf.Session():
  result = output.eval()
result
```

```
array([[2., 3., 4.],
       [5., 6., 7.]], dtype=float32)
```

Figure 5.4: Google Collaboratory uses Jupyter extensions to customize Jupyter for their users. The run/play icon to the left of the code cell is created using extensions. This is not present in the standard Jupyter software. TensorFlow is a library for creating Machine Learning experiments in Python.

Pandas (https://pandas.pydata.org/) provides resources for data analysis and more flexible data structures than core python. A good resource for getting familiar with these libraries is the Scipy Lecture Notes (https://www.scipy-lectures.org).

5.1.7   Advanced topic: extensions

There are many community contributed extensions that add functionality to Jupyter notebooks. Extensions vary from displaying an automated table of contents for a notebook, or prettify code, or hiding/showing solution cells. Here is the link for how to install and enable extensions: https://jupyter-contrib-nbextensions.readthedocs.io/en/latest/install.html

Here is a list of a collection of extensions that are bundled together: https://jupyter-contrib-nbextensions.readthedocs.io/en/latest/nbextensions.html

Creating custom extensions is a way to extend or customize Jupyter to add a capability that is not currently available with current extensions or out of the box. These extensions may be targeted for a specific kernel. Here are instructions for how to create and install custom extensions: https://jupyter-notebook.readthedocs.io/en/stable/extending/frontend_extensions.html

Figure X shows shows how Google Collaboratory, one of many tools to interact with Jupyter notebooks, leverages the power of Jupyter extensions for custom interaction and presentation.

5.2 Tips and tricks

5.2.1 Reminders

If you are using a single notebook as a standalone exercise in a traditional class (i.e., this is the only computational component of your class), then it is helpful to have a few cells at the top of that notebook that reviews how to navigate through the notebook and how to insert cells, etc.

5.2.2 Feedback

How do we get feedback from students in an interactive session to see if students have completed an exercise?

A low tech solution is to give students sticky notes of different colors, one meaning "finished" and one meaning "need help", that they can stick on the back of their computers. The instructor can then quickly look up to take a survey of the state of the class and decide how to proceed.

Projecting Slack or a similar chat group on a screen and having student copy-paste solutions (provided they are short functions) is a nice way to let everyone in the class see one another's solutions. A positive aspect of having multiple student solutions projected is that it can show the variety of ways to solve a problem. This gives an opportunity to talk about the readability of solutions and their efficiency. A downside is that in a large class, the shear volume of posts can make it overwhelming. Instead polling can be used to aggregate student answers and provide some form of feedback to the instructor. Nbgrader or travis-CI can also be options here, requiring students to submit completed code where it is assessed automatically. These will however require more setup and can take some time to complete.

5.2.3 Explaining each cell

A good rule of thumb is to always include a markdown cell above every code cell.

5.2.4 Custom styling

New notebook creators often try to centrally manage the formatting of headings, equations, and other textual items. For example, rather than using a standard markdown heading, a creator may over-design the headings by using HTML styles. This may create two problems:

1. The rendering of the notebook markdown may change and your formatted HTML header may not maintain the same look over time.

2. Headers created using markdown can be used by notebook tools, such as automatically creating a Table of Contents.

Our recommendation is to resist the desire to customize the styling and simply use the default representations. If you want to do customization (for example if you want to color certain cells) you can use CSS.

5.2.5 Length of notebooks

Notebook authors sometimes make the notebooks very long with many topics and sections. Notebook sections and cells are currently not easily reused in a copy/paste sense for mixing intra-notebook content. Until this functionality is available, we recommend that authors make short, self-contained notebooks around short topics. This allows other notebooks authors to mix and match notebooks to create curriculum.

Figure 5.5: Jupyter notebook displayed in a window pane inside Spyder.

5.3   Gotchas

5.3.1   Programming language   Jupyter

Teaching a class entirely with Jupyter can give the sense to students that this is the way all computational exploration is done. In particular, students can be confused into thinking that programming requires the notebook, instead of understanding that a notebook is just one way to interact with a particular language. This point should be made clear periodically. A good way to reinforce this is to show how to take a function that has been developed and debugged in a notebook and cut-paste it into a script (such as a file ending in .py for Python) and then import it into the notebook to regain that functionality. Also, the Integrated Development Environment (IDE), Spyder, has a plugin (https://github.com/spyder-ide/spyder-notebook) that allows notebooks to be displayed alongside Python scripts and a python terminal which can be useful for showing this dichotomy.

5.3.2   Restart, restart, restart...

Often, students may need to stop a computation, and this can be accomplished by pressing the "Interrupt" button in the toolbar. However, students should also be made aware of how to restart the kernel in a notebook, and what this means. There are several instances when students might need to do this. Sometimes students write code that can go into an infinite loop. The visual cues that notebooks give in this case are subtle, and students may not realize this and don't understand why the notebook is non-responsive. In live-coding situations, it can be useful to demonstrate this to students and show them how to restart the kernel and carry on.

A second instance of where restarting a kernel might be needed is due to how the notebook stores the state

of the computation. We like to think that, since the notebook is laid out in a linear fashion, that the state will always reflect what would happen if the notebook was run from the start up to that point. However, it is common to work in a notebook out of order, for instance if students ask a question about some previous example. If the variable has been changed in subsequent cells, then its value might not reflect what you expect when you rerun a cell earlier in the notebook. Restarting the kernel is sometimes the only solution.

5.3.3 Notebook hygiene

Many gotchas can be mitigated by developing notebooks that will be robust to incremental and non-linear execution. The main principle is to minimize side-effects of executing a cell and manifests itself somewhat differently in different languages; our suggestions here will be relevant to Python and may need to be adapted for other languages. Notebooks should generally be able to execute sequentially, such as via "restart kernel and run all cells". (An exception is when a notebook is intentionally incomplete for the purpose of live coding or student exercises, see nbgrader or the exercise estnations for more elegant ways to handle this.) Variable mutation is the most common way in which a notebook may malfunction when executing cells in a non-linear way (e.g., in response to student questions or when comparing and contrasting different methodologies). Sometimes this mutation is incidental, through dummy variables that were not meant to have significance outside the scope of the cell in which they are used. Their scope can be limited by placing them in a function, even if that function is only called once. Redefinition of functions can often be avoided by parameterizing the desired functionality as would typically be done if designing a library (though this may be a distracting software design for novice programmers). Function definitions should have little or no dependency on variables from their enclosing scope. When modifying cells for demos and formative assessments during class, it is useful to either copy the cell or modify/execute such that a conforming implementation remains present when moving on to other cells where it may be used. Additionally, you can minimize these issues by grouping code in a single cell that should always be executed sequentially, because code within a cell will always be sequential.

Chapter 6

Getting your class going with Jupyter

You have several options on how to get Jupyter notebooks to your students. You can ask students to install Jupyter on their own computers, install Jupyter on lab computers for students to use, or run Jupyter on a remote server that your students access on the internet.

6.1 Local installation on students' or lab computers

"Local installation" means that each computer is running the software that includes the Jupyter Notebook. Typically, this requires installing a distribution that includes Jupyter, Python, and possibly other language kernels.

A popular distribution for Jupyter is Anaconda, which is easy to install on Windows, Mac, and Linux. Because it can install everything at the user level, it does not require the installer to have administrator (or root) permissions.

> **What is Anaconda?** You will see the Anaconda distribution recommended by many educators and course authors. Anaconda is a package manager, an environment manager, a Python distribution, a collection of over 1,500+ open source packages, and also Jupyter. It is free to download, open source, and easy to install, on any computer system (Windows, Mac OS X or Linux). It also includes the conda packaging utility to update and install new packages of the Python and R ecosystems, and to manage computational environments. According to the company's webpage, Anaconda has more than 6 million users: https://www.anaconda.com/what-is-anaconda/. The Software Carpentry project provides installation instructions with videos for anaconda: https://swcarpentry.github.io/python-novice-gapminder/setup/.

Two other popular open source projects that can run Jupyter notebooks are nteract and Hydrogen. nteract is installed by downloading a binary installer from their website and double-clicking the installation file. Once nteract is installed, any Jupyter notebook on a student's local system can be double-clicked and it will open within nteract. nteract's simple user interface make it an excellent choice for students new to computer programming. Hydrogen is a very popular plugin for the open source Atom editor; it's currently used by over 700,000 people. Hydrogen lets a user edit, display, and execute a notebook within the Atom editor.

You can ask students to install Jupyter on their own computer or make it possible for them to use lab computers. These can also be combined: give students the instructions to install it on their own, but also tell them that it's available in the lab if they can't get it to work on their laptop. This way you don't need a large enough computer lab for everyone, and don't need to worry that not everyone can get it to work on their own.

6.1.1   Jupyter on student-owned computers

The benefits of installation on student-owned computers include:

- Once students have the software on their computers, they always have access to it; they can work anywhere, and they can use it for internships, jobs, and other non-school activities.

- It is easy for them to install additional packages later.

- Students learn to install and set up Jupyter, and software in general, which is a skill they are likely to need.

- The total computing power for the class scales with the number of students, as long as each student has enough CPU power and memory to support the intended applications.

- You can adopt Jupyter without support or resources from your institution.

- Students learn to use Jupyter on their preferred OS, e.g. Linux, Mac, or Windows, which means they are already familiar with the basic idioms of their OS.

Drawbacks include:

- This approach is only possible if every student owns a computer with enough capacity.

- Students with less powerful computers might be at an unfair disadvantage.

- Although installation is generally easy, it still takes time. The time you spend at the beginning of a class can be worthwhile for a semester-long course that uses Jupyter throughout, but it is a barrier to using Jupyter for a single module or one-off assignment in a course about something else.

- Also, the amount of time spent debugging esoteric problems scales with the number of students: a class of 25 students is bound to have a few people with 32-bit processors, incompatible libraries, out-of-date operating systems, over-zealous virus checkers, etc., and a class with 100 students will have four times as many. One work-around is to have students work in pairs: the probability that more than half of the students cannot get it working is reduced.

- Discrepancies in installed library versions can cause issues for students and may lead to different behaviors when students run code.

Although Jupyter is cross-platform and ideally behaves the same on Windows, Mac, or Linux, and distributions such as Anaconda also behave very similarly on all platforms, the instructions for installing and launching it are slightly different on each operating system, so fine-grained instructions such as "double click here" or "type this command" need different versions for Linux, Mac, and Windows users, which can be challenging when the instructor presenting the material has only one platform at their disposal. It is worth developing detailed instructions that the students can go through at their own pace, rather than relying only on a live demo in class that will only apply to a fraction of the students.

6.1.2   Jupyter on lab computers

Using lab computers instead of student-owned computers has the benefits of uniformity and improved equity. Each student will have exactly the same setup, and the instructions will work the same for everyone. This reduces the amount of individual tech support required and guarantees that all students have access to enough computational power.

However, this deployment has some disadvantages:

- Depending on how much control you have of the computer lab, you might need institutional permission and support.

- Students might be limited to working on assignments only when they are on campus and when computer labs are open, which might be an unfair disadvantage for non-resident students or those with full time jobs.

- It might be difficult to install additional packages as the need arises, and students might not be allowed to install packages they need for projects.

- Even in a computer lab, it can be difficult to maintain consistency across machines, and to keep all installations functional.

6.2   Jupyter on remote servers

Even when Jupyter runs locally, it runs as a web application; that is, it runs in a browser connected to a server. In a local installation, the browser and the server run on the same machine. But it is also possible to run the server remotely.

In that case, students don't have to install anything; they only have to run a browser and load a URL.

There are several ways to run Jupyter on a remote server:

1. You can run Jupyter on a server owned by you or your institution.
2. You can run Jupyter in a temporary environment running in the cloud.
3. You can run Jupyter in a persistent environment running in the cloud.

Running Jupyter remotely has many of the advantages of running in a lab: you can provide a consistent environment and guarantee that all students have access to sufficient computation resources. And it mitigates one of the drawbacks of a lab installation, since students have access to cloud resources from anywhere, not just on campus.

Working in the cloud also means that students do not have to manage their own backups of a laptop hard drive. Although a student could still inadvertently overwrite, delete, or destroy the contents of a notebook stored in the cloud, they will not lose their entire work if a laptop is damaged or lost.

For simple, one-off uses of Jupyter (say, for a single assignment or in-class activity) the cloud option is very attractive as it requires little in-class time to discuss installation of additional software.

6.2.1   Running in a temporary environment in the cloud

The easiest option for running Jupyter in the cloud is to use a cloud service that provides temporary environments. Some of these services are free of cost, and you can use them without installing anything.

These environments are well-suited for short examples in classes that do not use Jupyter extensively. Students can open a notebook and start running with the push of a button.

However, there are some limitations to these services:

- If your notebooks depend on particular packages, or particular versions of packages, it can be difficult to satisfy these requirements.

- These services run notebooks in a temporary environment that disappears if it is left idle. So they might not be suitable for managing student work.

- Some of these services do not guarantee a level of service and may not be as reliable as you need for a class or workshop.

  Binder (mybinder.org) is an open-source service provided by Project Jupyter. It allows the owner of a set of notebooks residing in a git repository to pre-build an image in the Binder service, and get a shareable link that any visitor can click to obtain a working instance of JupyterHub, with

the notebooks in the repository. The session is temporary (any changes the user makes will be deleted when closing the tab or window), but it's fully interactive. Binder is currently one of the favorite services for running one-off workshops or tutorials.

6.2.2 Running on servers you control

If you have access to a server or cluster with enough computing power to support your class—including CPU and especially memory—you can provide a Jupyter as a service using JupyterHub.

JupyterHub is open-source software that provides a cloud-based Jupyter application for each user in a group. Each user has their own account and home directory on the server. The Hub, JupyterHub's central system, allows authenticating users and starting individual Jupyter notebook servers. Programs that start notebook servers can use a variety of technical solutions. For more details, see https://github.com/jupyterhub/jupyterhub/wiki/Spawners

Once the Hub starts a user's notebook server, the Jupyter Notebook running in the cloud behaves just like Jupyter does when installed on an individual's computer, but JupyterHub will be running notebooks and storing files on a remote cloud computer. Students can download notebooks stored in the cloud to their local computer if they wish to work with a local installation as well. Additionally, students can upload notebooks (and other files) from their local computer to the cloud.

While anyone can run a JupyterHub server on their own Linux or Mac computer, installing and configuring JupyterHub requires sophisticated knowledge spanning the Linux/Unix operating system, system administration, and networking. For more information, see:

- https://github.com/jupyterhub/jupyterhub (the basic JupyterHub project, which can be installed on a bare-metal server, a virtual private server (VPS), or a commercial cloud cluster)

- https://github.com/jupyterhub/the-littlest-jupyterhub (a simplified installation of JupyterHub on a remote server or VPS)

- https://github.com/jupyterhub/zero-to-jupyterhub-k8s (a step-by-step guide to install JupyterHub on a Kubernetes cloud system)

Providing a JupyterHub service offers several benefits. First, students get up and running immediately—they spend no time installing software. They navigate to a web URL, log in to JupyterHub, and begin using Jupyter. This ability to quickly log in and begin computing is a powerful way to get students to engage with the lesson, builds confidence, and avoids the sometimes-stressful experience of installing software on the student's computer.

However, running JupyterHub on your own server has drawbacks:

- Getting started is not easy; most instructors would require (or at least benefit from) institutional support that may not be available.

- It can be difficult to scale: if the number of students increases, you might need more computing power. And the load students generate can be uneven; for example, if everyone runs a computationally-intensive example at the same time, your server might not be able to handle it.

- This option can be expensive, unless you already have servers with sufficient power.

6.2.3 Running Jupyter in the cloud

If you or your institution don't own computing hardware with the power to support your class, you can run JupyterHub on virtual servers provided by cloud services like AWS and Microsoft Azure. In those environment, you can install JupyterHub as described in the previous section

Commercial offerings also exist to use Jupyter in the cloud, some of which provide free trials or a "freemium" pricing model. They include:

- CoCalc, previously MathSageCloud (https://cocalc.com) is a subscription service with a free trial plan. The service also includes the ability to share files with project collaborators. The no-cost version does not allow network access. This has some important limitations (for example, you cannot install additional packages or kernels).

- Gryd (https://gryd.us) is another subscription service with a free tier. It includes course-management features, like a way to create a course, invite students, and deploy auto-graded assignments.

- codio (https://codio.com/features/ide)

- Microsoft Azure notebooks (https://notebooks.azure.com/)

- Amazon Sagemaker (https://docs.aws.amazon.com/sagemaker/latest/dg/ex1-prepare.html)

- Gradient by Paperspace (https://www.paperspace.com/gradient)

- Google Colaboratory (https://colab.research.google.com/)

The biggest advantage of these services is that they require no installation and minimal setup by instructors, and some of them provide features that integrate with learning management systems. However, instructors generally have to create student accounts and set up student environments.

These services are highly scalable; that is, they can handle large numbers of students and uneven loads. However, they are not infallible; they might require some tending to make sure students have access to enough resources.

The biggest drawback of these services is that they can be expensive. Some charge on a per-student basis, with limits on computation and memory use. Some charge on the basis of actual use, which can be unpredictable (and might require instructors to enforce limits on student activities).

Other drawbacks include:

- It may be difficult or impossible to install packages you, or particular versions of packages.

- Some of these services impose limits on what students can do; for example, they might have limited ability to access external services.

- Many of these services are relatively new, and they sometimes expose instructors and students to rough edges.

- Students generally lose access to their accounts when the class ends (or a limited time after).

- There may be privacy concerns with sharing student information on commercial servers. Some institutions have agreements with one or more of these providers that address privacy.

6.3 Distribution and collection of materials

You may want to distribute course materials to and collect them from students. A variety of options are available. Some important things to consider:

- Do you want to share your notebooks publicly, or do they require privacy?

- Can the notebooks that the students create or edit be public? Or do they require privacy?

- How do you plan to assess collected notebooks?

- Do you need integration with your LMS?

- Do you need integration with a file-sharing system?

- Do you want to distribute with the cell output showing?

- Do students need software that is not easily available on their own (or laboratory) computers?

Jupyter notebooks are plain text computer files, so you can distribute them to students and collect them using any system that handles text files, including GitHub, Google Drive, and (as a last resort) email.

6.3.1   Learning management systems

Many instructors use a Learning Management System (LMS) to communicate with students. These tools offer private file sharing and assignments that connect to the students' institutional computing accounts and they can be used to distribute and collect notebooks as text files. However, most LMS tools are not yet notebook-aware, so they don't render notebooks or make it easy for instructors to comment on or grade them.

Some tools and workflows are being actively developed to connect the Jupyter ecosystem to the LMS ecosystem using the Learning Tools Interoperability (LTI) standard. By the time you read this, you might find that the options have improved.

6.3.2   Web hosting

Notebooks can be publicly hosted on any website, so students can download the files by clicking on a link. Most web hosting software is not notebook-aware, but `nbviewer` is a web service that renders notebooks (https://nbviewer.jupyter.org/). Also, some browser extensions can open notebook files via `nbviewer` (https://github.com/jiffyclub/open-in-nbviewer).

GitHub Pages (and other similar services) can be used to host rendered notebooks, and continuous integration services can build the web pages from the notebooks and then display content. See Jupyter Book (https://github.com/choldgraf/jupyter-book) and use of `doctr` to do this.

6.3.3   Git

One of the most popular tools for distributing and collecting notebooks is Git, which is a version control system. Files under Git control are often hosted on services like GitHub, GitLab, and Bitbucket. Many of these services are notebook-aware; for example, when you view a notebook on GitHub, you see a rendered notebook that includes formatted text, typeset mathematics, code highlighting, and the output of the code, including figures.

Educators at academic institutions can use GitHub Classroom, which allows instructors to set up assignments for a class. Students click on a link for an assignment and a copy of the assignment repo is created and initialized with the assignment content, which can be a notebook. Each student's repository can be made private, with access only granted for the student and instructor. This can be an efficient way to distribute assignments to a large class.

A drawback of Git is that it is hard to use. It might be worth spending time in your class to teach Git if it is valuable for students to learn about version control. But if this is not one of the learning goals for your class, you can minimize the students' exposure to Git using graphical interfaces like GitHub Desktop and Git for Windows.

The default Git tools for comparing files and merging changes do not work well with Jupyter notebooks. However, there are specialized tools to help with these tasks (see Notebooks Under Version Control).

6.3.4 JupyterHub

If your students are using JupyterHub, you can place notebooks and any related files directly into the students' directories manually or via a script. If `nbgrader` is available on your JupyterHub instance you can use it to collect and distribute notebooks (whether or not you choose to use `nbgrader`'s assessment features). This allows you to develop the notebooks and incrementally make them visible to the students for them to "fetch". They can then edit the notebooks or create new ones in the directory created in their storage space, and then publish their notebooks back to you for downloading, viewing, or assessing with the `nbgrader` tools (see the next section for details on this tool).

> **What is nbgrader?** `nbgrader` is a tool for creating, handling, and automatically grading assignments based on Jupyter notebooks. It works as a Jupyter extension that the course creator installs on their computer. `nbgrader` is a flexible project in the Jupyter ecosystem that allows the distribution and collection of materials. As its name implies, it also can grade assignments; it can be used in a distributed manner where each student is running Jupyter on their own computers, or in a centralized manner, for example, if the students each have an account on a JupyterHub installation. (More details in the Assessment section.) https://nbgrader.readthedocs.io

6.3.5 Using an LMS and `nbgrader` together:

Integration of `nbgrader` with learning management systems is still primitive, but the following is a strategy that works with current tools.

1. The instructor creates an assignment notebook using `nbgrader`, then distributes the assignment to students via an LMS.

2. Students complete the assignment and upload the solution to the LMS.

3. The instructor downloads the completed assignments as zip file and extracts the students' solutions in a Jupyter environment.

4. Instructors and graders use `nbgrader` to grade the assignment and save the grades to a CSV file.

5. The CSV file is then uploaded to the LMS.

There are already some tools that make this workflow easier, including the Extractor plugin to the ZipCollect feature in `nbgrader` (https://nbgrader.readthedocs.io/en/stable/plugins/zipcollect-plugin.html).

6.4 Assessing student learning with Jupyter notebooks

Many educators develop course-assessment activities as Jupyter notebooks. This includes exams, in-class activities, homework assignments, and projects.

One simple way to handle the assessment of a notebook-based submission is to have students either print them out, email them, submit them as a standard electronic document (say into a Course Management System), or drop them into a shared folder. At that point, the instructor can mark and grade them in a traditional manner, for example by simply writing comments on a printout or adding annotations to a PDF.

```
**_Pro Tip: Printing out a notebook can sometimes result in wasted
  space on pages, especially for notebooks with many images or
  figures. Converting to PDF requires large/complex LaTeX
  installations. Exporting to HTML and then printing often gives a
  better result._**
```

`nbgrader` allows code cells in a notebook to be marked to be auto-graded or manually graded. An instructor can then create an assignment that can be completely auto-graded, requiring little work after the notebook

has been created. This makes grading much easier and scales well with large class sizes. However, creating such an auto-graded notebook in `nbgrader` can be quite time-consuming. In addition, pedagogically a completely auto-graded notebook may have serious downsides. For example, studies suggest that students learn better when they can actively connect a topic to their own interests [CITATION NEEDED]. One method of encouraging this is to have a "reflection" question on each submission. Such a reflection question can encourage students to comment on the material in a personal way, but it cannot be auto-graded. Another downside is that simply autograding code with unit tests is unlikely to assess many of the learning objectives you might have for an assignment, e.g., ability to use specific software-design patterns. To address this, you can create manually graded cells for a portion of an assignment and provide written feedback to the student.

GOTCHA: At the time of this writing, nbgrader has some limitations that require careful use. For example, using it in a multi-class setting (say, on JupyterHub) requires that instructors coordinate the naming of assignments so that they do not collide.

`nbgrader` is a sophisticated tool that can be set up to allow multiple graders, teaching assistants, and more. For more information on using `nbgrader`, see https://github.com/jupyter/nbgrader.

There are some third-party notebook-based assessment solutions. For example CoCalc (<www.cocalc.com>) and Vocareum (<www.vocareum.com>) provide a cloud notebook platform that can also perform assessment similar to nbgrader.

_For example, cocalc.com offers... [are there other third-party course management notebook-oriented solutions?] and Berkeley uses DataHub for their large Data8 course. Vocareum (https://www.vocareum.com) TODO

6.5   How do you create Jupyter notebooks for reuse and sharing?

As you create notebooks for your lectures, computational essays, or homework assignments, you may wish to think about how to make it possible that they can be reused by yourself and others.

First, you may want to make the materials openly accessible and findable via the internet. This suggests avoiding keeping the notebooks behind a "walled garden," such as a Course Management System. That is, users may have access to some material, but be prevented from seeing other materials. You will have to decide whether you want others to have full access. For example, many teachers do not want students to be able to see notebooks that may have solutions, or hints of solutions and therefore limit their access.

> To share your notebook with others you can submit it to https://www.engage-csedu.org/. This curated collection of open educational computing resources is maintained by the National Center for Women in Information Technology (NCWIT).

If you decide to make your notebooks reusable by others, make it clear under which license the materials can be used. For example, you can include a Creative Commons attribution and share-alike statement at the bottom of your notebook. Adding a license allows people to reuse your materials without asking for permission explicitly.

GitHub may be the most common service to host and share notebooks, where they can be viewed (including rendering), downloaded, or forked by others. (Private repositories can also be used to limit visibility to colleagues, students, or other organizations.). Make sure to be aware of some of the pitfalls of keeping notebooks version controlled however (see Notebooks Under Version Control for details).

Another potential issue with sharing deals with external files that you may want to include in your notebook. This is in contrast with content (say a plot) that can be directly created by the notebook's code. Possible content includes data, images and videos using code and embedding tags in markdown or HTML. The implication is that if you share your notebook you must include the external files along with the notebook. This can be done a number of ways including using a version control repository, a zip-file, or a file sharing service. Another external dependency issue with sharing notebooks involves software libraries. In this case

you share a configuration file that a user can use to setup the same environment. Examples of these files include a conda `env.yml`, a pip `requirements.txt`, or `dockerfile`.

Because Jupyter notebooks embed the output of cells into the ipynb file itself (e.g., images, videos, etc.), the files can grow large. To make it possible to display the cell output via the renderers on Github, Gitlab, or nbviewer, save the notebook after it is executed and then upload to those services. If instead you want to reduce file size and provide the notebooks to someone with the code cell output cleared, choose this option in the notebook's dropdown interface. Then the user will need to execute the notebook themselves to see the output.

6.6 Jupyter: a 21st Century genre of Open Educational Resources and practices

Educators create teaching and learning materials. With the appearance of the internet, a community of educators began producing open access traditional teaching materials. In parallel, a community of software developers began creating open source software. Each community developed their own development patterns. In particular open source software communities gravitated to the bazaar style[1] of distributed and collaborative work. Jupyter notebooks may be the first time that these communities are merging. Jupyter notebook authors are applying the content creation patterns they use to the creation of open educational resources that teach computation or teach through computation.

Open Education encompasses a large community, with its own conferences and journals, with leaders and advocated practices. The most visible efforts are related to Open Educational Resources (OER): the creation and adoption of openly licensed learning materials. In 1994, Wayne Hodgins coined the term "learning object" and the idea spread that digital materials could be designed and made to be *reused*. This was followed by efforts to develop metadata standards, content exchanges, and so on (addressing the concern of how to find the objects to reuse them). In 1998, David Wiley coined the term "open content" and spread the idea that principles of Free and Open Source Software (FOSS) could be applied to content on the World Wide Web (OpenContent, 1998). The Creative Commons non-profit organization was founded in 2001 to provide ready-made license agreements for sharing content and served a vital infrastructure role on the spread of OER. The Creative Commons licenses are now the most widely used licensing framework for open education. The year 2001 also saw the launch of MIT OpenCourseWare (OCW). MIT promised free public access for non-commercial uses of their course materials. It was a unique commitment at an institutional level, strengthened by the MIT brand. Other universities joined the OCW movement: Rice with the OpenStax project (now formerly Connexion), CMU with the Open Learning Initiative (OLI), Utah State University with the Center for Open and Sustainable Learning, and so on. Today, the Open Education Consortium has hundreds of members from around the world. The recurrent topics in OER are: reducing costs for students buying textbooks, increasing access, and dealing with copyright and licenses.

In the last few years, educators using Jupyter have been creating and sharing all kinds of educational materials in the form of notebooks, typically under a Creative Commons Attribution license (CC-BY). In fact, Jupyter is *a new genre of OER*. But in addition to creating open content, educators using Jupyter often take active part in the Jupyter community and adopt the *culture* of open-source software. This is a culture with strong ethical commitments, related to freedom of access, transparency, and governance (Coleman, 2012). The content they create has the value of giving access (the very definition of OER), under an open model. But open-source culture also promotes a culture of collaboration. In this regard, engaging in teaching with Jupyter opens new possibilities for educators to engage in *open development* and collaborate with others in producing lessons, tutorials, courses, and even books.

---

[1]The bazaar style is a method of collectively creating software that isn't top down directed like a traditional company hierarchy.

Chapter 7

Usage case studies

Contributors to this chapter: you may increase adoption by new users if you integrate information about some of the following into your case:

1. Demonstrate that you can increase students' ability to:
    1. Engage material & participate in class
    2. Understand material and perform well
    3. prep their for career
    4. Enjoy learning this way
2. describe:
    1. how it fits with how their students learn
    2. how it connects to how they teach
    3. the needed resources (support, hardware, etc.)
    4. the necessary logistics (e.g., how much time will it take? Be honest: time is a consideration, and an important reason that people do not adopt new practices, but is not a reason that they stopped using one)
    5. what Jupyter does in terms of promoting learning, instructor affordances

7.1 Jupyter notebooks in support of scaling for large enrollments

7.1.1 Supporting large enrollment courses at UC Berkeley

The University of California at Berkeley started a pilot course titled "Foundations in Data Science" (also known as Data-8) for about 100 incoming undergraduate students in Fall 2015. Data-8, the fastest growing course in Berkeley's history, is entirely Jupyter-based, allowing the program to scale the course to 1,400 students in 2018. This scale is made possible by Jupyter's shared computational environment. In particular, Jupyter allowed "browser-based computation, avoiding the need for students to install software, transfer files, or update libraries" (see The Course of the Future and the Technology Behind It [https://data.berkeley.edu/news/coursefuture]). Data-8 is powered by JupyterHub and all the course materials are published openly (http://data8.org).

7.1.2 Large-scale adoption: Jupyter across Canada

Recognizing the importance of data science, computational research, and educational resources, the Pacific Institute for the Mathematical Sciences (PIMS), in partnership with Compute Canada and Cybera, have launched JupyterHub platforms (under the project name Syzygy) to support researchers and educators

across Canada. Syzygy (http://syzygy.ca) provides access to cloud-hosted Jupyter resources using existing institutional credentials and encourages the development of computational and data science skills. It is currently accessible at 16 institutions across the country (McMaster, Queen's, SFU, UAlberta, UBC, UCalgary, ULethbridge, UNewBrunswick, UOttawa, URegina, USask, UToronto, UVic, UWashington (US), UWaterloo, Yorku) and has been used by over 11,000 people at those institutions.

Syzygy is extensively used for teaching, but is also being used for research activities. One notable example is a scientific software seminar at the University of British Columbia, where graduate students and post-doctoral researchers meet to share and learn data science techniques with their peers. Initiatives are also underway, as part of syzygy, to deepen its relevance into research by providing seamless access to larger and more varied types of resources (GPUs, parallel machines, different language kernels etc.).

Callysto (https://callysto.ca/) is a related project, also launched by PIMS and Cybera, to bring Jupyter to students in Canadian middle and high schools (grades 5-12). Callysto focuses on creating and curating open content (https://github.com/callysto). This content forms the basis of project workshops, where teachers can work through the materials interactively, before taking them back to their classrooms. The content links to a supporting JupyterHub installation (integrated with the authentication systems for the networks of school districts) allowing easy access to the materials and a Jupyter environment to learn and create in.

– Ian Allison

7.1.3 Quick switch: moving an existing course to Python and Jupyter (at the last minute)

For many years, our chemical engineering kinetics course had used software for differential equation and nonlinear simultaneous equation solving to simulate reactors and solve design problems. The software, recommended and described by the textbook, was installed in the college's computer labs, but licenses for student-owned computers were expensive and it was only available for Windows. In Spring 2015, I was informed my class now had 52 students, but the largest computer lab had room for only 40. As the semester progressed and we neared the chapters that required numerical simulation, I rewrote the examples using Python and SciPy and created Jupyter notebooks, walking students through the steps involved in setting up and solving the problems. I found Lorena Barba's open-source MOOC materials online, and adapted these for my "getting started" notebooks. I had students install Anaconda on their own computers, and got everyone up and running without any central infrastructure or support from the college's IT staff. I found the Jupyter Notebook format of including "lecture note" style commentary along with short, unintimidating, snippets of code, to be extremely effective. A couple of years later I passed on the course to a new instructor, who took my course materials, taught himself some Python, and continued to use Jupyter notebooks for content delivery and assignments.

The first year was a bit rough around the edges as I introduced it quite late in the semester. Still, it is clear that the approach resonated with students. An alumnus from my 2016 course wrote, "I thought that your course was very successful, especially the use of Jupyter Notebook as a classroom and assignment tool. I still remember specific problems that we went over in class (e.g., the microfluidic reactor array with heterogeneous catalysis), and I feel that the use of Python to solve problems throughout the course greatly benefited my understanding of fundamental concepts. I went on to use Python [in the pharmaceutical industry], where I built tools for bioinformatics data analysis, mutation network profiling for protein engineering experiments, and RNA structure prediction from experimental data and molecular thermodynamics."

– Richard West

7.2   The "CFD Python" story: guiding learners at their own pace

"CFD Python" is a collection of Jupyter notebooks based on a practical module that I began using in class in my Computational Fluid Dynamics (CFD) course at Boston University in 2009. The 5-week module

develops worked examples that build on each other to incrementally guide the learner to create a program to solve the Navier-Stokes equations of fluid mechanics, in 12 steps. In 2013, I was invited to teach a mini-course in the Latin-American School in High-Performance Computing, in Argentina. The Jupyter notebooks platform allowed me to create a guided narrative to support learners with different background experience and knowledge. For that event, we wrote notebooks based on the CFD course module, to use as instructional scaffolding in the 2-full-days of minicourse. Twenty students worked through the notebooks as self-paced lessons, while I went from desk to desk asking and answering questions. About four of the students completed all the lessons in the 2 days, a bulk of them achieved up to about Step 8, and a few of them lagged behind in Steps 4 or 5 by the end of the course. For those who completed the full module, they had achieved in 2 days what my regular students in the classroom normally took 5 weeks to do. Seeing that was an eye-opening moment: both the power of worked examples in code, and the ability to allow learners to follow their own pace made a remarkable difference in these learners.

REF — Barba, Lorena A., and Forsyth, Gilbert F. (2018). CFD Python: the 12 steps to Navier–Stokes equations. Journal of Open Source Education, 1(9), 21, https://doi.org/10.21105/jose.0002

Based on the experience developing the "CFD Python" learning module, we adopted this basic design pattern for creating lessons using computable content:

1. Break it down into small steps
2. Chunk small steps into bigger steps
3. Add narrative and connect
4. Link out to documentation
5. Interleave easy exercises
6. Spice with challenge questions/tasks
7. Publish openly online

– Lorena A. Barba

7.3   Analyzing music with music21

I became interested in learning more about Python in 2013 after reading a tutorial by Luciano Ramalho as he was writing Fluent Python. Since I tend to seek out projects that match my outside interests (music, art, and nature) I was looking for Python projects with music and came across Myke Cuthbert's music21 project. Music21, an open source music theory and analysis library maintained by Professor Michael Cuthbert at MIT, provides a set of tools to answer questions about music quickly and simply. Users can create, analyze, and share music with just a few lines of code. Myke's use of the notebook hooked me. Unlike many things that I had worked on before, the notebooks made it easy to get started and to write small code snippets that did real work! The more I used the notebooks and showed them to people that I taught at Fab Lab San Diego, the more that I saw the power of the notebook to engage a user and empower them to explore and learn.

Music, a universal language, appeals to learners of all origins, ages, education levels, and interests. As a subject that casts a wide appeal, music offers the opportunity to engage and delight learners. It's an accessible subject that has a low barrier to entry for learners from disciplines beyond computer science and engineering.

– Carol Willing

Education benefits

- lessons notebooks can be tailored to age appropriate content within music
- multisensory
- ability in K12 to align with the standards
- possibilities for bringing in multi-subject learning
  - writing

- – history
- – math
- – science
- accessibility through audio and braille

Misc quotes (perhaps pick a couple?):

"I think of music21 as being composed of two parts. The first is infrastructure, routines for reading, writing, and manipulating musical scores, while the second consists of a higher-level analytical toolkit—generating a Roman numeral from a chord and key, putting chords into normal form, checking for parallel fifths, identifying scales containing a given pitch or chord, and so on." —Bruce Tymoczko, Professor of Music, Princeton

Inclusive

"It's not exclusive, but inclusive, which is the whole spirit of jazz." — Herbie Hancock

Education

"So, you can't stay in one place, no matter how comfortable that place is. It's all about growing."—Mavis Staples

Universal

"Music in the soul can be heard by the universe."—Lao Tzu

Communication

Music is the greatest communication in the world. Even if people don't understand the language that you're singing in, they still know good music when they hear it."—Lou Rawls

"In the beginner's mind there are many possibilities. In the expert's mind there are few." –Shunryu Suzuki

7.4 Interactivity in computer science (high school and middle school)

Who

High school and middle school students at Cal Poly SLO's EPIC program completed a two hour workshop on Interactivity in Computer Science. The workshop participants included dual language learners (English as a Second Language) and students who have had limited access to computers prior to the workshop.

Why

Providing early access to at-risk groups who may not see themselves as capable of learning to code or use computation

Illustrate that there are many skills beyond math and science that are needed to create software applications

What

Two hour workshop that maximizes "hands on" exploration with the goal of building an ongoing interest in computer science

- short lectures
- interactive discussion - LISTEN
- hands on - DO/APPLY This section is self-paced to engage different learning styles and prior knowledge
- recap - DISCUSS
- 8 or so projects with achievements outlined
- modern curriculum including p5.js, jupyter, binder, deep learning and machine learning with Tensor-Flow and Magenta (art and music)
- Goal is to empower students to understand that they CAN use CS to solve real world problems

Instructor Approach

- Start with high quality engaging content
- Self contained notebooks
- Use widgets to add additional interactivity

7.5 Interactive geophysics with Jupyter

The GeoSci.xyz project (https://geosci.xyz) is an effort to develop a community of scientists and educators around learning resources and software for the geosciences. The project includes multiple open-source textbooks, each which have associated Jupyter notebook "apps" that serve as interactive simulation engines for exploring concepts in geophysics. We have used these resources in an undergraduate course on applied geophysics at the University of British Columbia; this course is primarily taken by by geologists and engineers (non-geophysics majors). In 2017, we delivered a 2 day short course for professionals, graduate students, and researchers in 26 different countries around the world (https://disc2017.geosci.xyz). In both of these courses, the goal is to provide learners with an overview of the various geophysical methods (e.g. magnetics, gravity, seismic, electromagnetics) and concepts governing the physics; we do not dive into details of the math nor do we expect students to program or write any lines of code. The role of Jupyter notebooks in these courses is to serve as a tool for visualizing and exploring the physics.

During a lecture, the notebooks as a presentation medium lend to a dynamic presentation style, where we as instructors can select model parameters based on student input. Concepts are reinforced as students then use these same notebooks in labs and assignments. We have found that the notebook apps are most effective when students are first asked to critically think about what they expect to see and then visualize the result. If the resultant image matches their expectation, then they understand the concept, and if not, it is an opportunity to learn and further explore.

– Lindsey Heagy

7.6 Investigating hurricanes

Who

Middle school and high school students visiting Columbia's School of Engineering and Applied Sciences on a field trip

Why

Students often come through looking to tour labs and experience some of the research that is being done at the school. Unfortunately certain fields, in this case computational mathematics and hurricane research, do not lend themselves to these types of events.

What

Instead of a lab or lecture a computer lab was reserved for an hour and a Jupyter notebook used to walk students through some basic visualizations and data analysis encouraging students to change the code displayed to answer questions such as "Where did Hurricane Sandy go?" and "What storms occurred during 1981?". This includes a number of visualizations of hurricane tracks, coloring by strength of storms, and an analysis of average number of storms per year. Notebook is available at https://github.com/applied-math/demos.

– Kyle T. Mandli

Figure 7.1: Notebook "app" for exploring the direct current resistivity experiment over a two layer earth (https://em.geosci.xyz/apps.html).
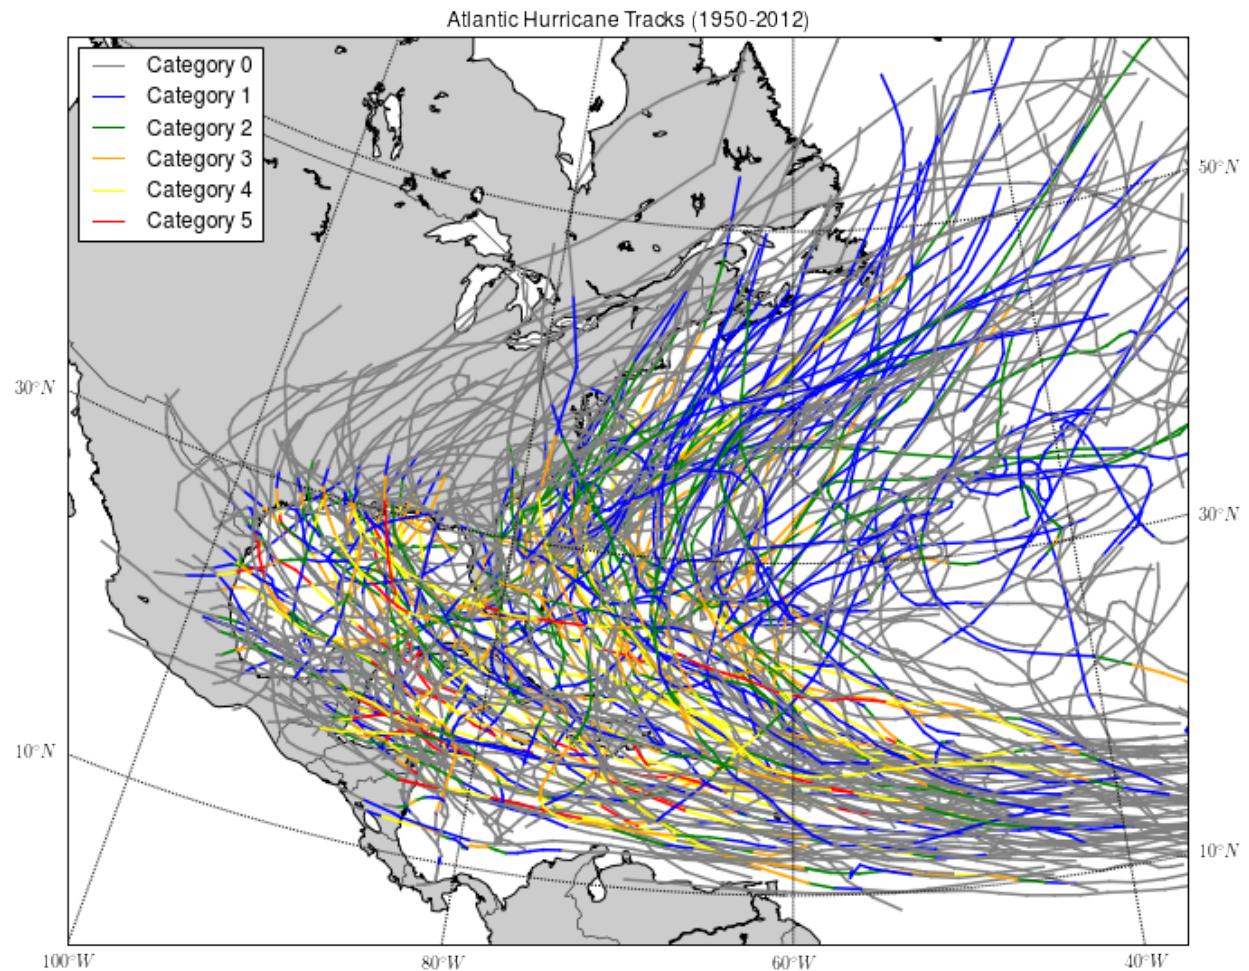
Figure 7.2: Visualization from the notebook at https://github.com/applied-math/demos demonstrating the paths of Atlantic hurricane tracks from 1950-2012 with coloring demonstrating category of storm.

Chapter 8

About the authors

8.1 Project lead

Lorena A. Barba

- George Washington University
- labarba@email.gwu.edu
- @LorenaABarba

Lorena A. Barba is Associate Professor of Mechanical and Aerospace Engineering at the George Washington University. She adopted Jupyter in 2013 and since then used it in every course she teaches. Her open course materials are well known and used by thousands of learners: CFD Python and Numerical MOOC are the best examples.

8.2 Authors at the sprint

Lecia J. Barker

- University of Colorado Boulder
- lecia.barker@colorado.edu
- @leciab

Lecia Barker is an Associate Professor and Associate Chair of Undergraduate Studies in the Department of Information Science at the University of Colorado Boulder. She is also a Senior Research Scientist for the National Center for Women & IT. Her research group is studying the diffusion and adoption of teaching practices in undergraduate computer science. Lecia holds a Ph.D. in Communication from CU Boulder and an MBA in Marketing from San Diego State University.

Douglas Blank

- Bryn Mawr College
- dblank@brynmawr.edu
- @dougblank

Douglas Blank is Associate Professor in the Department of Computer Science at Bryn Mawr College, a small, all-women's college outside of Philadelphia, PA, USA. He has a joint Ph.D. in Cognitive Science and

Computer Science from Indiana University, Bloomington. For over 20 years, Douglas has taught all levels of Computer Science. For the last 4 years, he has used Jupyter notebooks exclusively in the classroom. Douglas has published in the areas of Computer Science Education, Robotics, Artificial Intelligence, and Deep Learning. He is on the advisory board of Engage-CSEdu.org, a joint project between Google and the National Center for Women and Information Technology (NCWIT). Douglas also writes text and code at his website douglasblank.com.

Jed Brown

- University of Colorado Boulder
- jed@jedbrown.org
- @five9a2

Jed Brown is an Assistant Professor of Computer Science at the University of Colorado Boulder. He has been teaching numerical and scientific computing courses using Jupyter Notebook and `nbgrader` for three years, and leads a research group that develops computational methods and community software for computational science.

Allen Downey

- Olin College
- downey@allendowney.com
- @AllenDowney

Allen Downey is a professor of Computer Science at Olin College and the author of a series of open-source textbooks related to software and data science, including *Think Python*, *Think Bayes*, and *Think Complexity*, published by O'Reilly Media. These books, and the classes based on them, use Jupyter notebooks extensively. Prof Downey holds a Ph.D. in computer science from U.C. Berkeley, and M.S. and B.S. degrees from MIT.

Tim George

- Project Jupyter
- tgeorgeux@gmail.com

Timothy George is the Lead UI/UX Designer for Project Jupyter, focusing primarily on JupyterLab. In addition to his formal duties, Tim is also in working with Jupyter on design strategy, future products, governance, diversity and inclusion. He studied HCI at UC Irvine's Donald Bren School of Informatics and Computer Science where he received a Master's Degree.

Lindsey Heagy

- University of California Berkeley
- lindseyheagy@gmail.com
- @lindsey_jh

Lindsey Heagy is a Postdoctoral Researcher at the University of California Berkeley working on Project Jupyter and Jupyter in the geosciences. She recently completed her PhD at the University of British Columbia in geophysics. She is a project leader of GeoSci.xyz, an effort to build collaborative, interactive, web-based textbooks in the geosciences, and a core contributor to SimPEG, an open source framework for geophysical simulation and inversions. The GeoSci.xyz project relies heavily on Jupyter for making the content come to life.

Kyle Mandli

- Columbia University
- kyle.mandli@columbia.edu
- @KyleMandli

Kyle Mandli is an Assistant Professor in the Department of Applied Physics and Applied Mathematics at Columbia University. He has developed a set of openly available course notes centered around Jupyter notebooks and uses Jupyter for homework in conjunction with nbgrader. His other research interests include development of computational methods for coastal hazards such as storm surge and tsunamis.

Jason K. Moore

- University of California, Davis
- jkm@ucdavis.edu
- @moorepants

Jason K. Moore is an Assistant Teaching Professor of Mechanical and Aerospace Engineering at the University of California, Davis. He currently teaches dynamics and mechanical design related courses. He utilizes Jupyter notebooks to teach modeling and simulation and is working on a textbook about Mechanical Vibrations. He is responsible for the Jupyter related features in the LibreTexts project and is also a core developer of the SymPy and PyDy projects which utilizes Jupyter for training workshops, e.g. PyDy Tutorial and SymPy Code Generation Tutorial. Jason has PhD, MSc, and BSc degrees in mechanical engineering from UC Davis and Old Dominion University.

David Lippert

- Leidos

David Lippert is a software engineer at Leidos in Arlington, Virginia. He utilizes Jupyter notebooks primarily for exploratory data analysis and for training and evaluating machine learning algorithms. He has written Jupyter notebooks to create new Dr. Seuss sonnets and to evaluate if the Rotten Tomatoes Tomatometer can be trusted. He has a BA in computer science from Middlebury College.

Kyle E. Niemeyer

- Oregon State University
- kyle.niemeyer@oregonstate.edu
- @kyleniemeyer

Kyle Niemeyer is an Assistant Professor of Mechanical Engineering in the School of Mechanical, Industrial, and Manufacturing Engineering at Oregon State University. He teaches courses in numerical and analytical methods for solving differential equations as well as gas dynamics, and recently developed a graduate course on software development for engineering research. His research group develops and applies methods for modeling combustion and chemically reacting fluid flows. He is also on the steering committee of the Cantera open-source project for chemical kinetics, thermodynamics, and transport processes.

Ryan Watkins

- George Washington University
- rwatkins@gwu.edu
- @parsingscience

Ryan Watkins is a Professor of Educational Technology at George Washington University in Washington DC. He leads the Human-Technology Collaboration (HTC) PhD program area, and he teaches courses in needs assessment, instructional design, and research methods. Ryan's research focuses on how people and organizations define and assess needs. He is co-host of Parsing Science, a podcast where researchers share the stories behind their science. He also developed the We Share Science platform for sharing video abstracts of research.

Richard West

- Northeastern University
- R.West@northeastern.edu
- @richardhwest

Richard West is Associate Professor of Chemical Engineering at Northeastern University in Boston. He leads a research group in computational modeling for complex reacting systems like combustion or catalysis. He is a core member of the Cantera open-source project. As well as in an elective on "computational modeling in chemical engineering", he has integrated Python and Jupyter into core classes on chemical kinetics and reactor design, at both the undergraduate and graduate levels. As part of his NSF CAREER award, he is developing modules to teach students to use Python and SciPy to solve chemical engineering problems.

Elizabeth Wickes

- University of Illinois at Urbana-Champaign
- wickes1@illinois.edu
- @elliewix

Elizabeth Wickes is a Lecturer at the School of Information Sciences at the University of Illinois at Urbana-Champaign. She teaches foundational programming from an information and data sciences perspective, as well as other coursework on open data and reproducibility. Her programming course lectures are written in Jupyter notebooks and the class is taught via live coding.

Carol Willing

- Cal Poly San Luis Obispo
- willingc@gmail.com
- @WillingCarol

Carol Willing is a Research Software Engineer at Cal Poly San Luis Obispo working full-time on Project Jupyter. She is a Python Software Foundation Fellow and former Director; a Project Jupyter Steering Council member; and a core developer on CPython and Jupyter. Carol has an M.S. in Management from MIT and a B.S.E. in Electrical Engineering from Duke.

Michael Zingale

- Stony Brook University
- Michael.Zingale@stonybrook.edu
- @Michael_Zingale

Michael Zingale is an Associate Professor and computational astrophysicist at Stony Brook University. He has a PhD from University of Chicago (2000). He frequently teaches numerical methods and Python for scientific computing graduate courses, relying on Jupyter notebooks and python for much of the presentation. He is an advocate for open educational resources, as a founder of the Open Astrophysics Bookshelf project where he hosts his *Introduction to Computational Astrophysical Hydrodynamics* text.

Chapter 9

Glossary

Anaconda: a free, open-source package manager, environment manager, Python distribution, and collection of over 1,500+ open source packages including and also Jupyter. https://www.anaconda.com/what-is-anaconda/

API (Application Programming Interface): a specification of what a programmer must write or define to interact with a software library.

Binder: a hosted service that allows anyone to launch their own sandboxed notebook environment from a Git repository. https://mybinder.org

cell: the area in a Jupyter notebook where you can enter markdown, or computer code.

cloud, in the: used to describe software or documents hosted on a remote computer accessed over the internet.

CSV (Comma Separated Values): referring to a comma-separated value file. A plain-text file format such that each line is a list of data separated by commas.

DataFrame A common tabular data structure with rows and columns available in R and in Python through Pandas.

execute: technical term for having the computer perform the instructions of your program. Alias for "run it."

extension, Jupyter: in this instance, it is not a request for more time. Rather, a Jupyter extension is a bit of code, often developed by a third-party, that adds additional functionality to Jupyter. For example, a popular extension is a Table of Contents creator.

flipped classroom: a teaching style where students work on their own outside of class to learn new material (sometimes by watching recorded lectures or reading descriptive/interactive notebooks) and the come together in the classroom to practice what they've learned through exercises or experiments.

Git: a popular version control system (VCS) used for keeping track of changes of files over time.

IDE (Integrated Development Environment): software that assists in the development of additional software.

Jupyter: The term "Jupyter" may refer to one of a couple of different things: a community of users and developers focused on the open source software; the collection of tools and standards that, together, allow projects like the Jupyter Notebook to operate. The name refers to the three core programming languages supported: Julia, Python, and R.

JupyterHub: a cloud service that can provide access to Jupyter notebooks and environments to multiple users via a modern web browser. http://jupyter.org/hub

kernel: In Jupyter, a kernel is the packaging up of a language, and related programs needed to run it. For example, Python2 and Python3 are separate kernels.

LMS (Learning Management System): a cloud service that helps instructors manage aspects of classrooms.

load: how many students can a computer support?

Markdown: a text format that allows for basic formatting (headers, text styles, links) mixed inline with the text. Markdown files usually have the extension `.md` and can be rendered natively by GitHub and other tools.

magic: a meta-command typically starting with one or two percent signs. Changes the meaning of the contents of a line (one percent sign, `%`) or the cell (two percent signs, `%%`) from code to a particular meta-instruction. For example, `%%R` indicates that the cell contents will be interpreted as commands to the R language. Magics are kernel-specific (e.g., vary with the kernel in use).

nbgrader: a tool for creating, handling, and automatically grading assignments based on Jupyter notebooks. https://nbgrader.readthedocs.io

nbviewer: a web application for rendering Jupyter notebooks as static web pages, providing a URL to share and view them with a modern web browser. https://nbviewer.jupyter.org

nbconvert: a tool for converting Jupyter notebooks into other formats such as PDF, HTML, LaTeX, Markdown, reStructuredText, and others. https://nbconvert.readthedocs.io

notebook hidden state: a technical term referring to the value of variables that may have surprising results due to cells having been executed in a non-sequential order.

open source: software and documents that are created in a manner that give you rights to be able to use, and reproduce.

pattern: A "pattern" is a technical term referring to an abstract description of a labeled process. For example, "wash, rinse, repeat" is a common pattern for cleaning various objects.

scaffold: A teaching and learning pattern that provides steps in the learning process that build on prior learned knowledge.

service, JupyterHub: JupyterHub can take advantage of additional separate, but integrated, software extensions. These are called "services."

script: a colloquial term for a computer program.

unit test: a technical term for a "test" for checking to see if software is operating correctly.

URL (Universal Resource Locator): the address of a resource (e.g., webpage) on the internet.

widget: a user interface (such as buttons, sliders, and checkboxes) that allow the easy control of hidden computer code.

References

Barba, L., & Forsyth, G. (2018). CFD Python: The 12 steps to Navier–Stokes equations. *Journal of Open Source Education*, *1*(9), 21. https://doi.org/10.21105/jose.00021

Chen, O., Kalyuga, S., & Sweller, J. (2015). The worked example effect, the generation effect, and element interactivity. *Journal of Educational Psychology*, *107*(3), 689. https://doi.org/10.1037/edu0000018

Coleman, E. G. (2012). *Coding freedom: The ethics and aesthetics of hacking.* Princeton University Press.

Freeman, S., Eddy, S. L., McDonough, M., Smith, M. K., Okoroafor, N., Jordt, H., & Wenderoth, M. P. (2014). Active learning increases student performance in science, engineering, and mathematics. *Proceedings of the National Academy of Sciences*, *111*(23), 8410–8415. https://doi.org/10.1073/pnas.1319030111

Haller, H., & Krauss, S. (2002). Misinterpretations of significance: A problem students share with their teachers. *Methods of Psychological Research*, *7*(1), 1–20. Retrieved from http://www.dgps.de/fachgruppen/methoden/mpr-online/issue16/art1/haller.pdf

Meurer, A., Smith, C. P., Paprocki, M., Čertík, O., Kirpichev, S. B., Rocklin, M., … Scopatz, A. (2017). SymPy: Symbolic computing in Python. *PeerJ Computer Science*, *3*, e103. https://doi.org/10.7717/peerj-cs.103

Moore, M. G. (1989). Editorial: Three types of interaction. *American Journal of Distance Education.* https://doi.org/10.1080/08923648909526659

OpenContent. (1998). About opencontent. Retrieved 18 December 2002 from http://opencontent.org/.

Raymond, E. S. (1996). *The new hacker's dictionary.* MIT Press.

Sweller, J. (2006). The worked example effect and human cognition. *Learning and Instruction*, *16*(2), 165–169. https://doi.org/10.1016/j.learninstruc.2006.02.005